

# Classification

Machine Learning Course - CS-433

Oct 19, 2021

Nicolas Flammarion

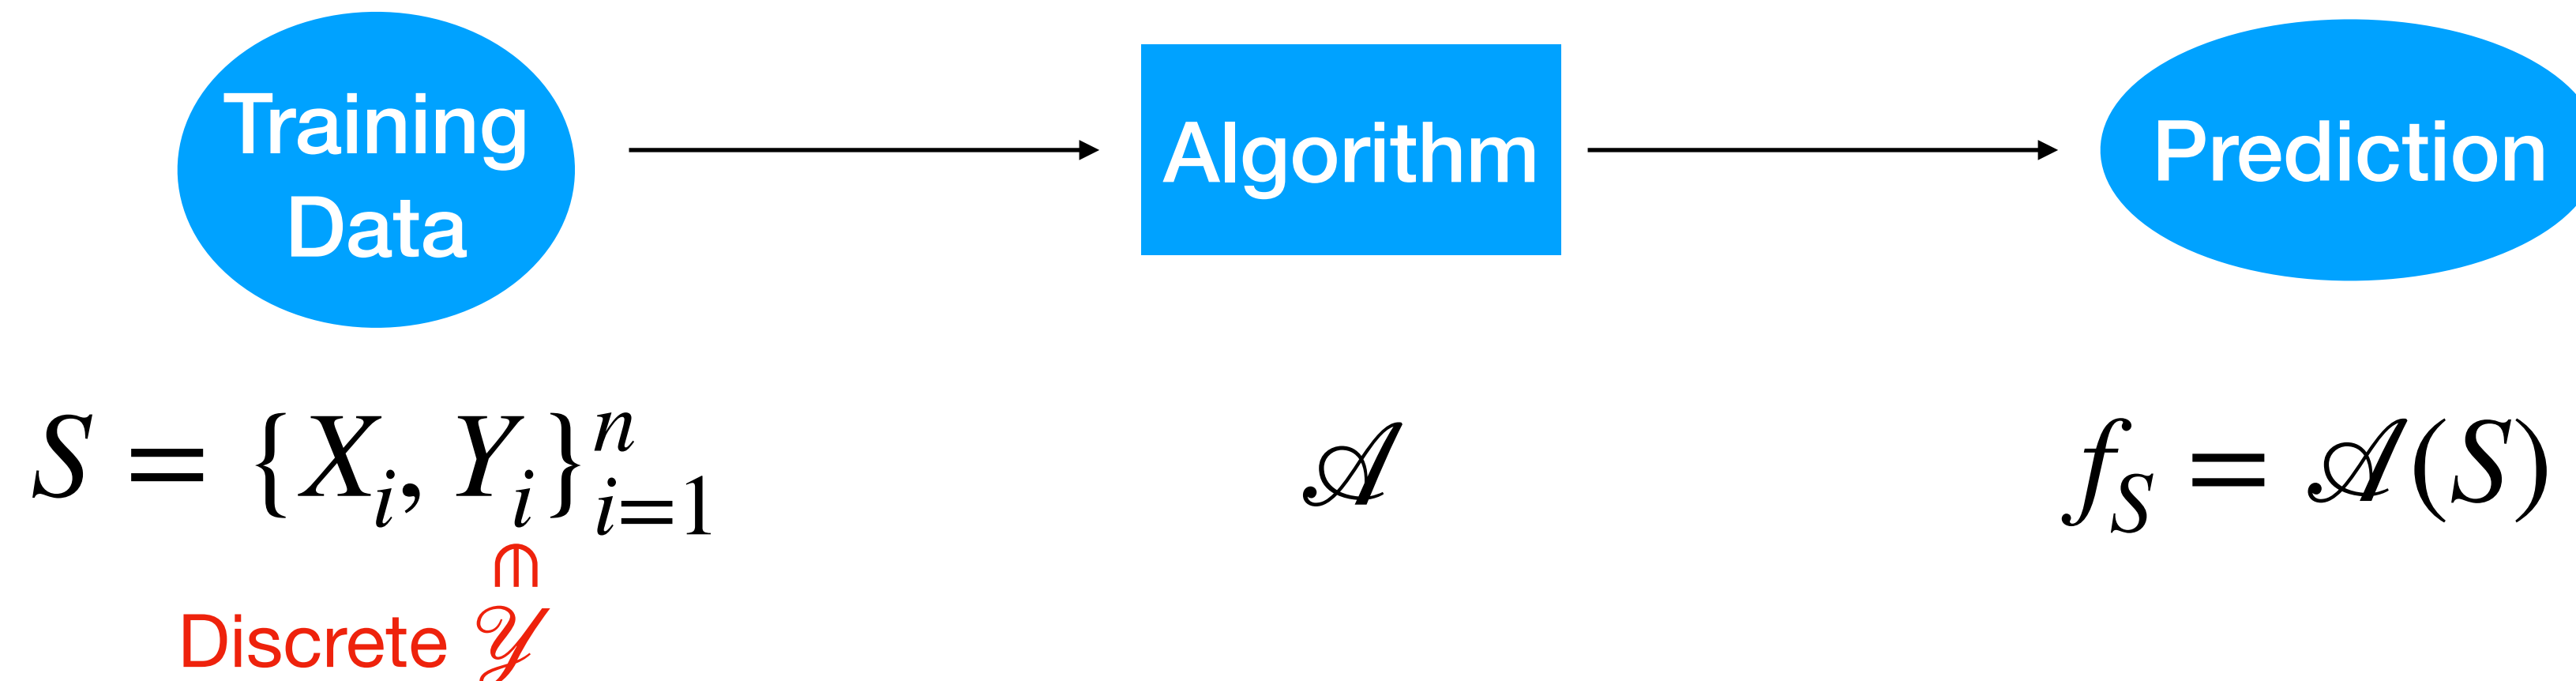
**EPFL**

# Definition of classification

We observe some data  $S = \{X_i, Y_i\}_{i=1}^n \in \mathcal{X} \times \underbrace{\mathcal{Y}}_{\text{Discrete Set}}$

Goal: given a new  $X$ , we want to predict its label  $Y$

How:



# Classification: relates input to a categorical variable

$$(X, Y) \in \mathcal{X} \times \underbrace{\mathcal{Y}}_{\text{Discrete Set}}$$

**Binary Classification:**  $Y$  can take two values

$Y \in \{c_1, c_2\}$  where  $c_i$  are the class labels . We often use  $\{0, 1\}$  or  $\{-1, 1\}$



**Multi-class classification:**  $Y$  can take more than two values

$Y \in \{c_1, \dots, c_{K-1}\}$  for a  $K$  class problem. We often use  $\{0, \dots, K - 1\}$



no ordering between classes

# Spam Detection



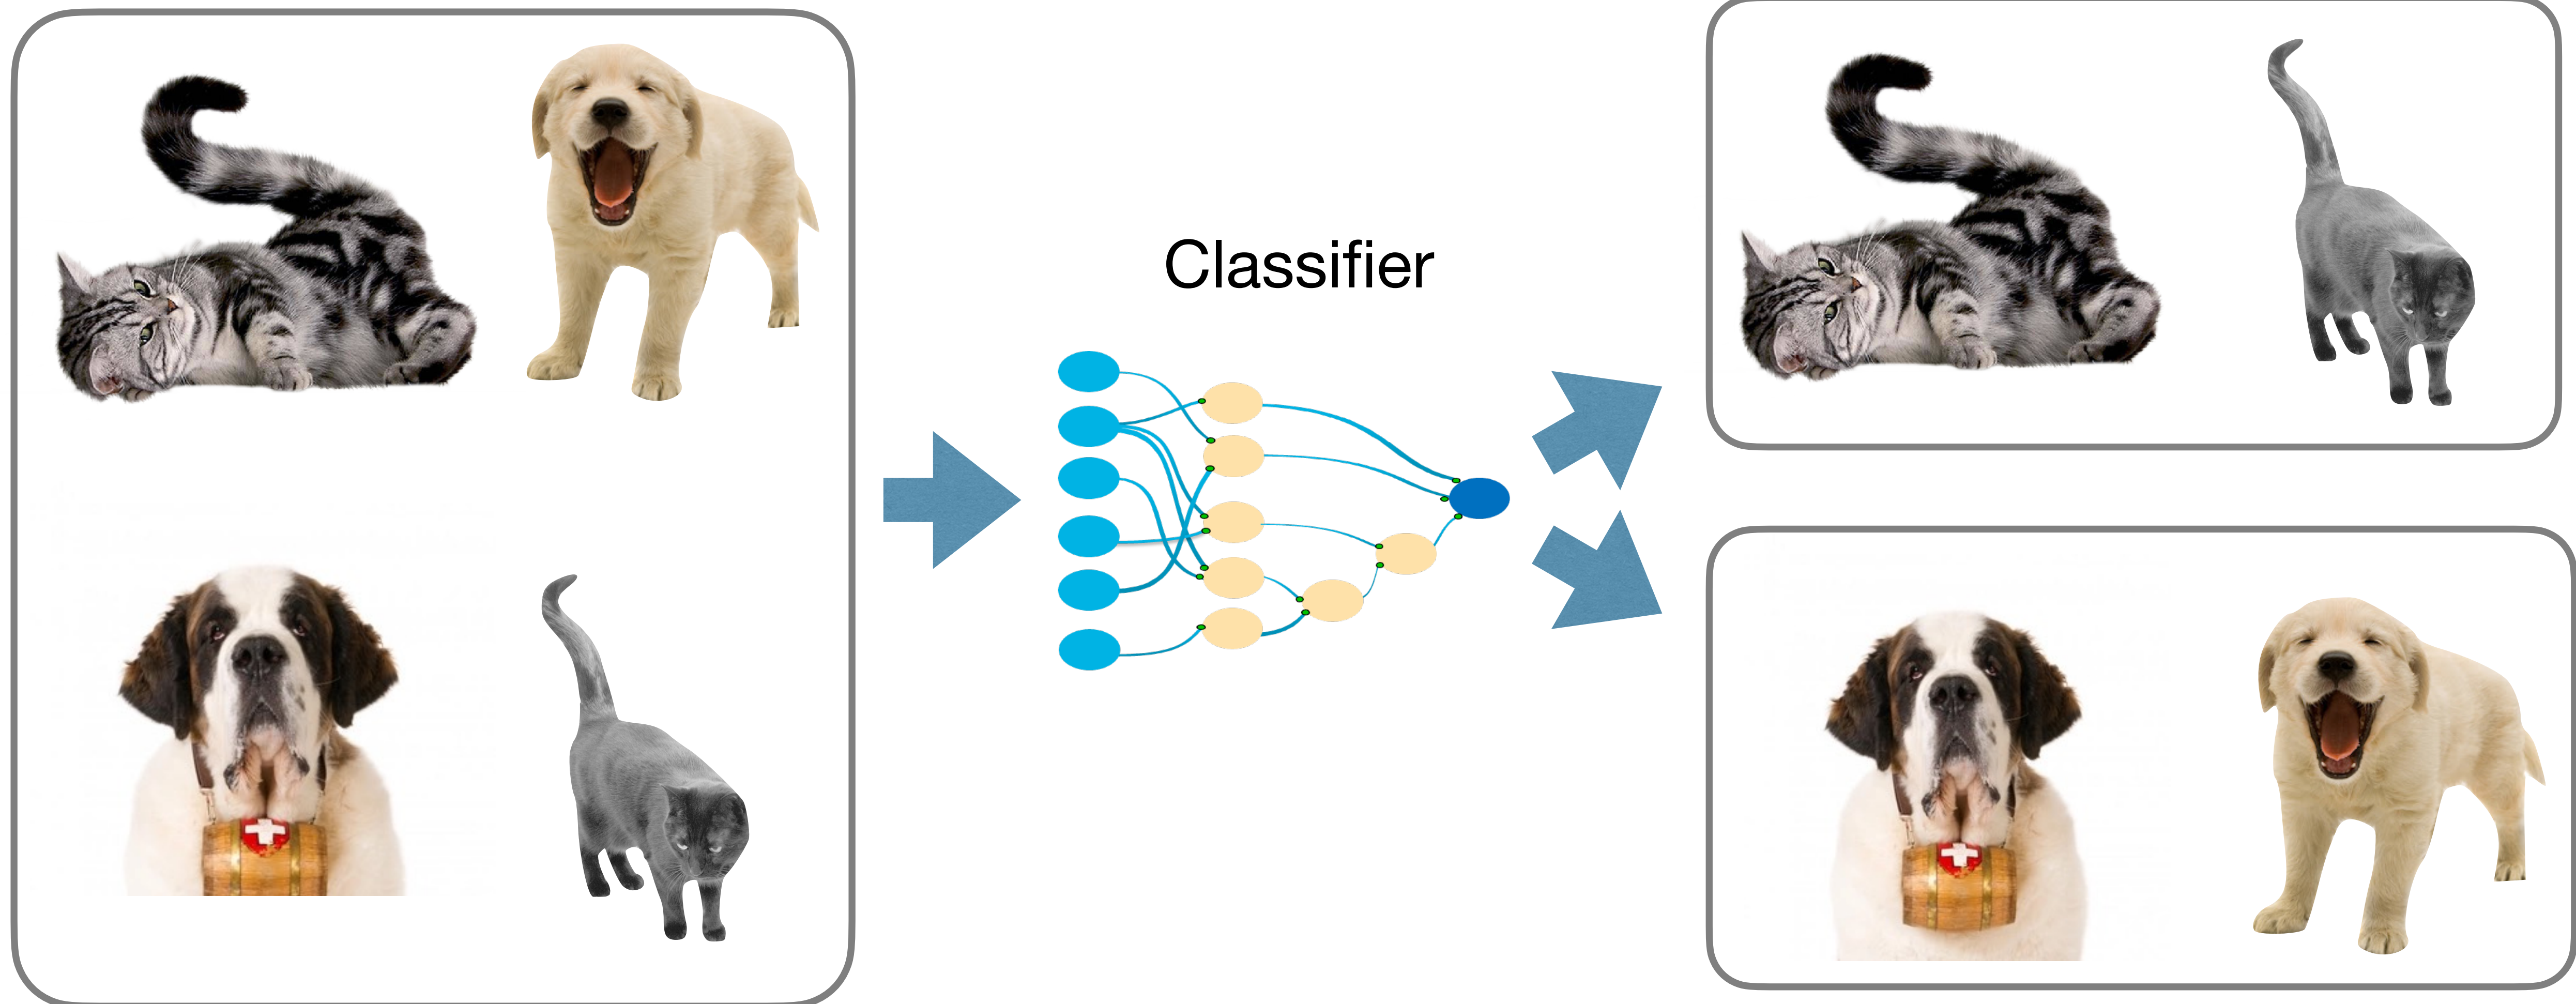
**Inbox**




**Spam Folder**



# Image classification



# Credit Card Default




- + individual who defaulted on their credit card payments
- individual who did not



# Classifier


A classifier  $f: \mathcal{X} \rightarrow \mathcal{Y}$  divides the input space into a collection of region belonging to each class



# Classifier

A classifier  $f: \mathcal{X} \rightarrow \mathcal{Y}$  divides the input space into a collection of region belonging to each class

It can be linear




Linear Decision boundary



# Classifier

A classifier  $f: \mathcal{X} \rightarrow \mathcal{Y}$  divides the input space into a collection of region belonging to each class

It can also be nonlinear



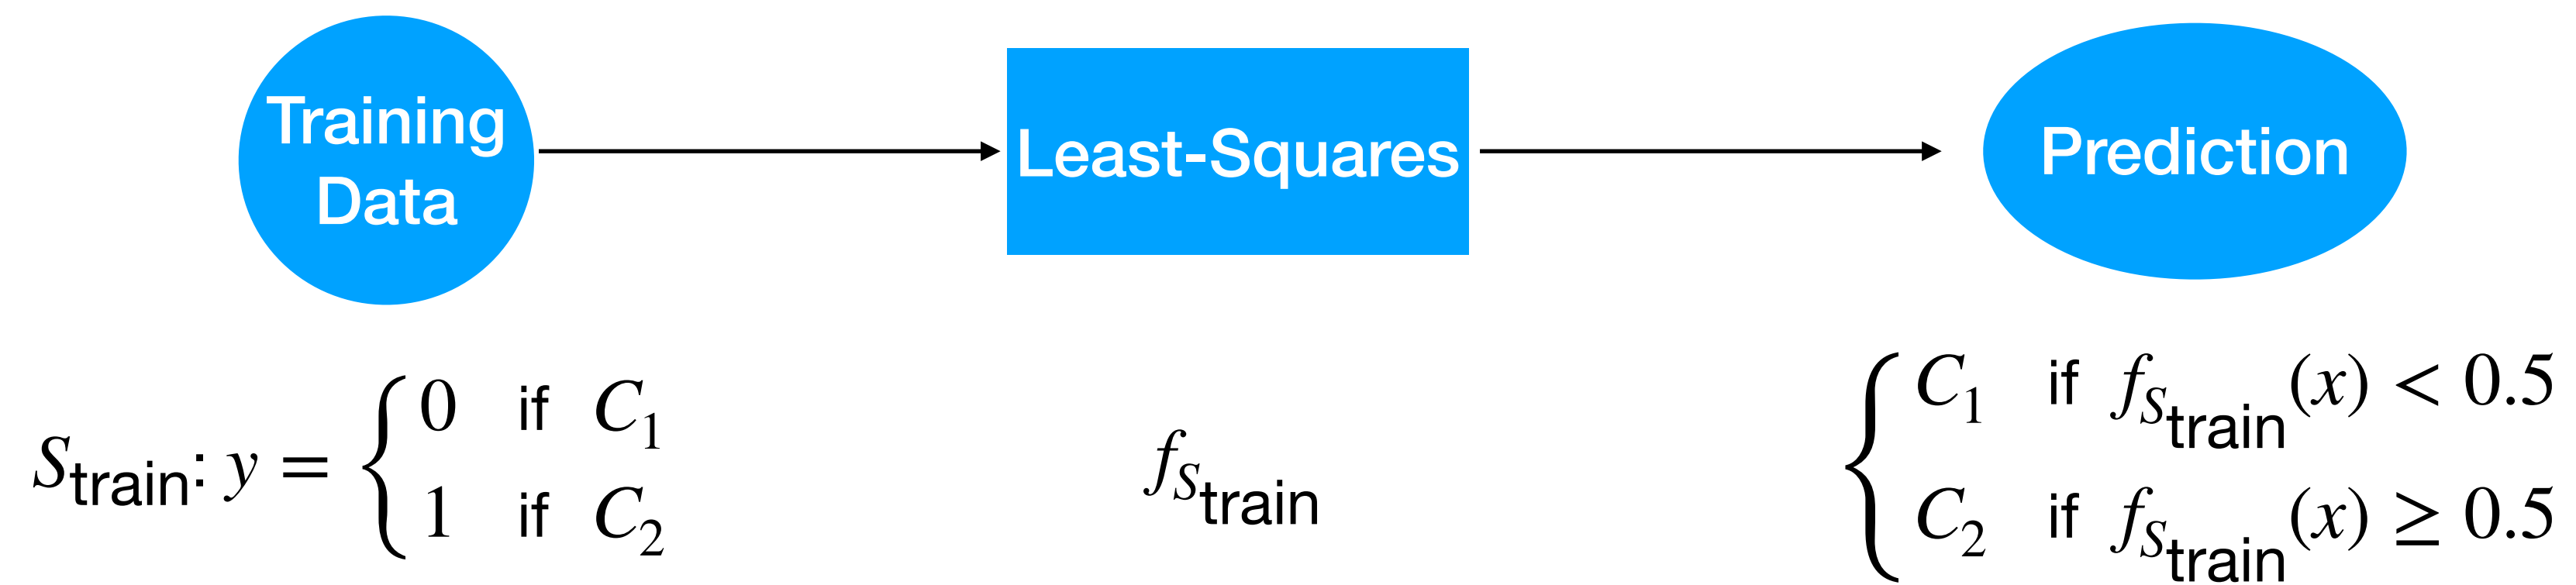
**Nonlinear Decision boundary**

# Classification: a special case of regression?

Classification is a **regression problem** with discrete labels:

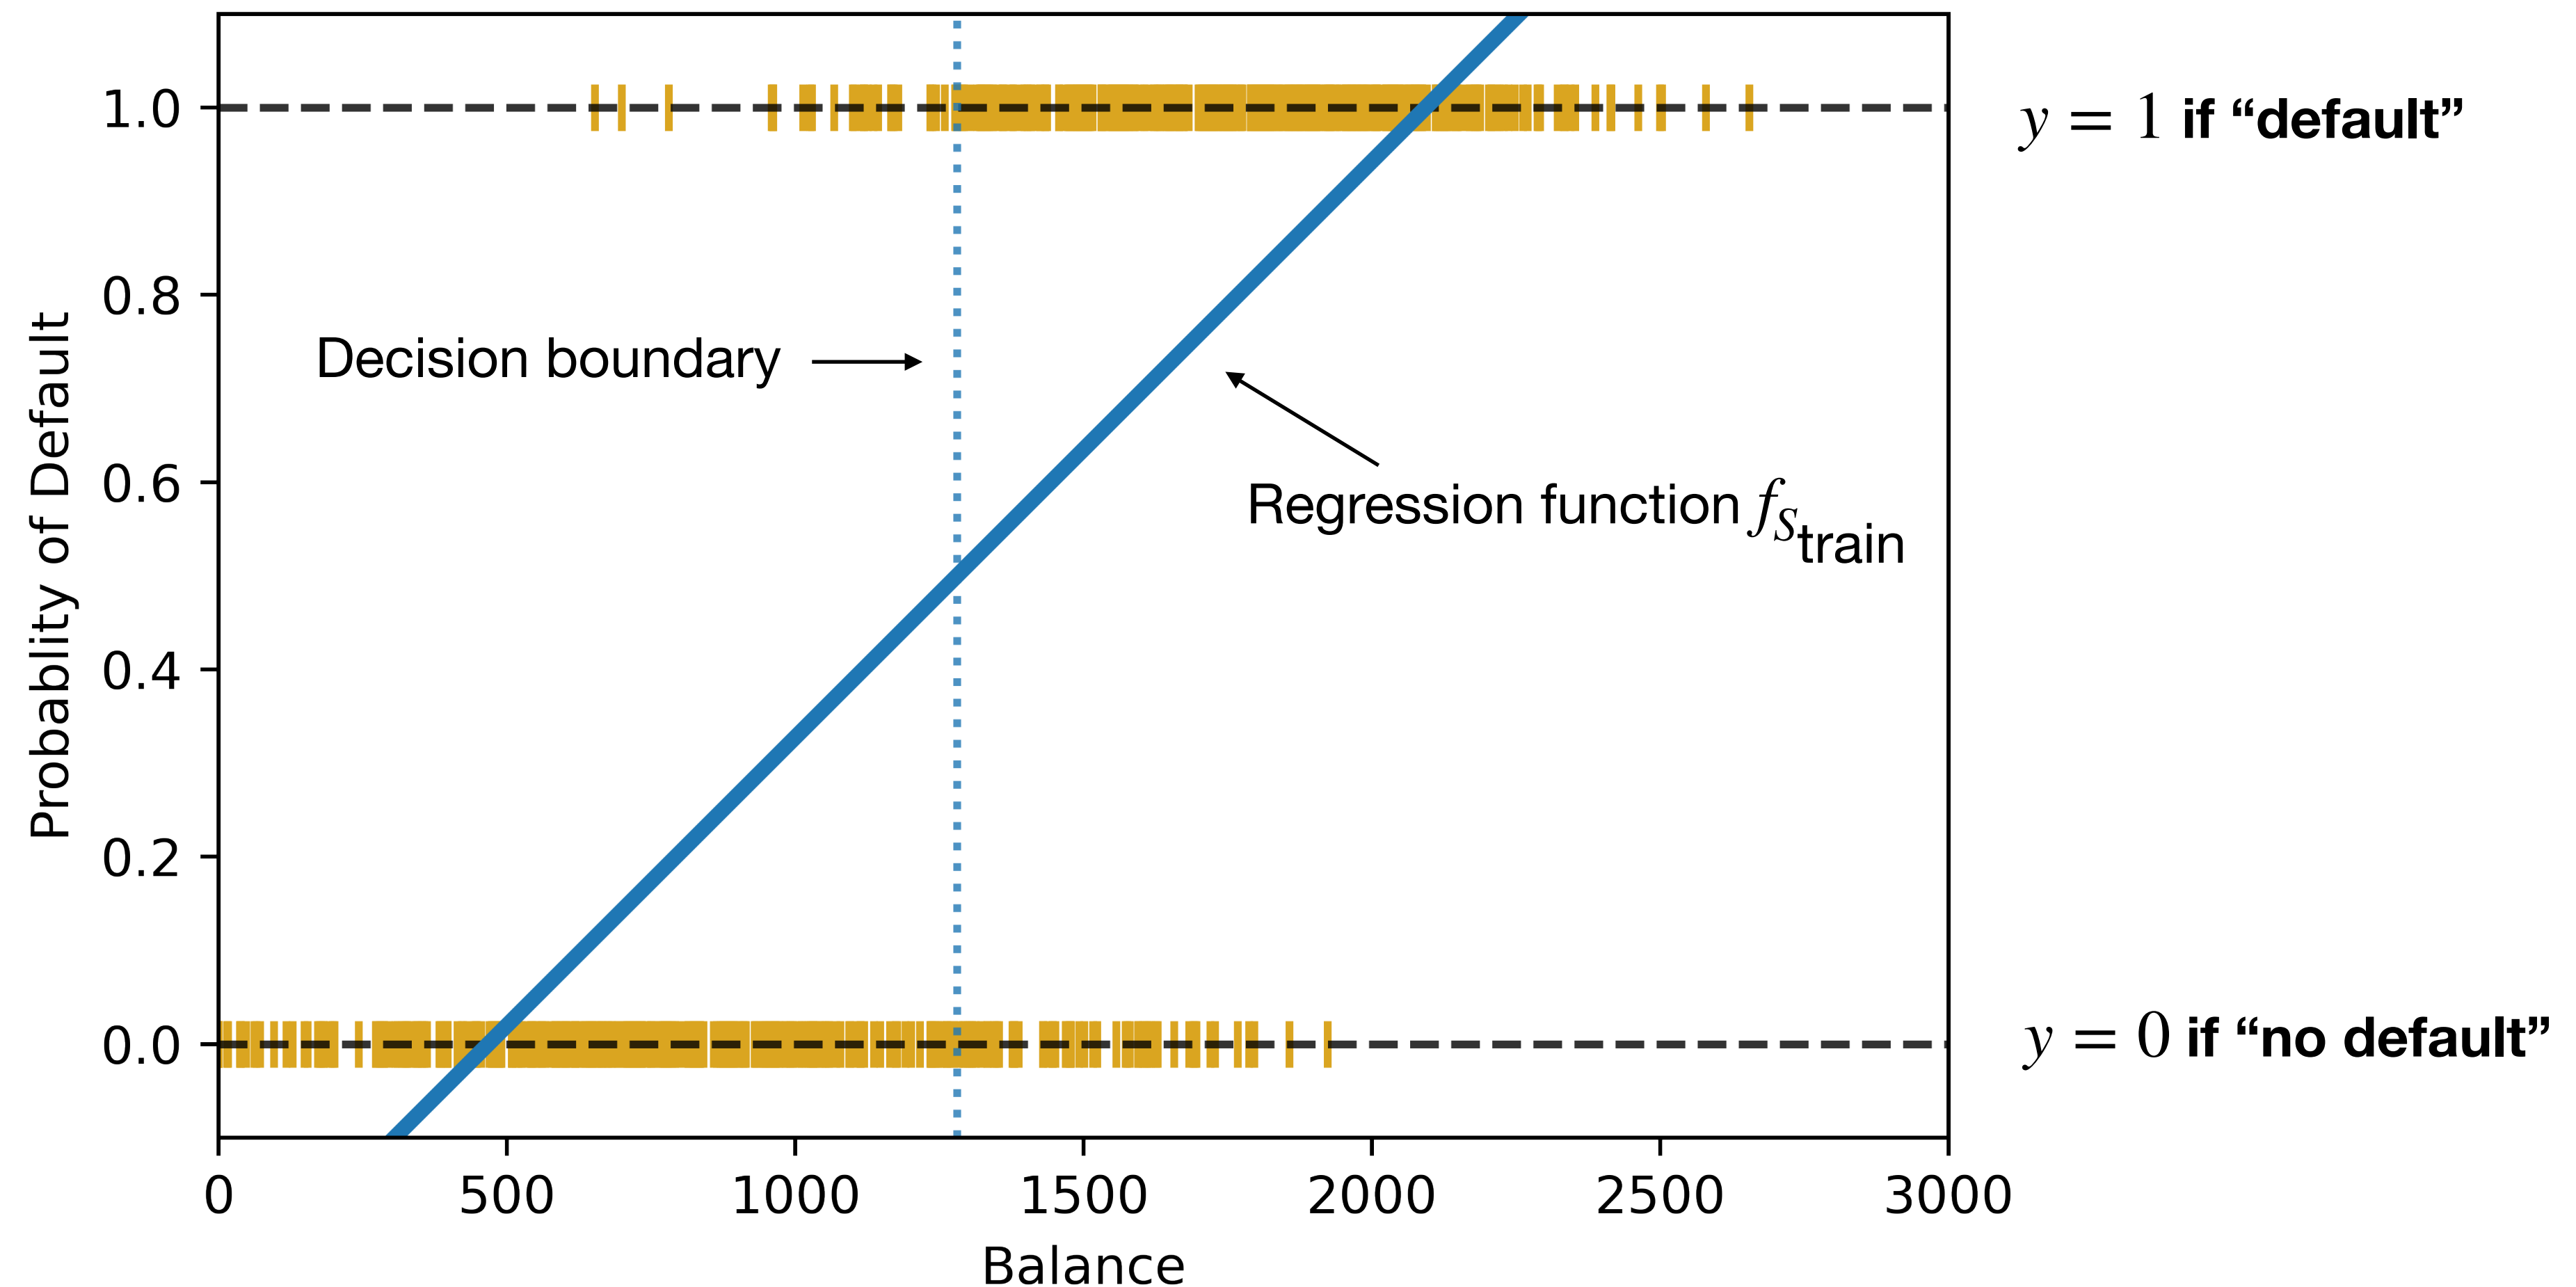
$$(x, y) \in \mathcal{X} \times \{0, 1\} \subset \mathcal{X} \times \mathbb{R}$$

Could we use previously seen regression methods to solve it?



# Is it a good idea?


Credit-card default problem:



We label the output as probability  
for sake of interpretation

# Classification is not just a special form of regression


A. The predicted values are not probabilities (not in  $[0,1]$ )





# Classification is not just a special form of regression



## B. Sensitivity to unbalanced data



The position of the line depends crucially on how many points are in each class

# Classification is not just a special form of regression

## C. Sensitivity to extreme values:



The position of the line depends crucially on where the points lie


**Why:** the square loss we used for regression is not suitable for classification

# How to perform classification?



- A lot of different approaches have been developed
- We will not detail them exhaustively today
- Rather we will provide quick introductions

- Fundamental task of classification:

**separate the space into various decision regions**




# Nearest Neighbor





# Nearest Neighbor

Assume that nearby points are likely to have similar label



You assign the label of the closest point in your training set.


# Nearest Neighbor

## Pros:

- **No optimization** or training
- **Easy** to implement
- Works well in **low dimension** where you can get some very complex decision boundaries


## Cons:

- **Slow** at query time
- Bad for high dimensional data
- Choice of **local distance** is crucial




# k-Nearest Neighbor

A new point  $X$  is classified by a **majority vote among the k-nearest neighbor of  $X$**



# k-Nearest Neighbor

A new point  $X$  is classified by a **majority vote among the k-nearest neighbor of  $X$**




Generalization: smoothing kernels; weighted linear combination of elements



# Linear Decision boundaries

Assume we restrict ourself to linear decision boundaries (hyperplane):


➡ Prediction:  $g(x) = \text{sign}(x^\top w)$



# Separating hyperplane

Assume we restrict ourself to linear decision boundaries (hyperplane):


Assume the data are linearly separable, i.e, it exists a separating hyperplane



Which separating hyperplane would you pick?

# Margin


Key concept: The margin is the distance between the hyperplane and the closest point



➡ Take the one with the largest margin!

# Max-margin separating hyperplane

Pick the hyperplane which maximizes the margin




Why: If we slightly change the training set, the number of misclassification will stay low  
➡ It will lead us to support vector machine (SVM) and logistic regression



# Non linear classifier

- Linear decision boundaries will not always work.
- Features augmentation ( $X, X^2, X^3, X^4$ )
- Kernel Method



**Do we still have time?**

# A little bit of theory

$$(X, Y) \sim \mathcal{D} \text{ with } X \in \mathcal{X}, Y \in \mathcal{Y} = \{0,1\}$$

Loss function:

$$\ell(y, y') = 1_{y \neq y'} = \begin{cases} 1 & \text{if } y \neq y' \\ 0 & \text{if } y = y' \end{cases}$$

True risk for the classification:

$$L_{\mathcal{D}}(g) = \mathbb{E}_{\mathcal{D}}[1_{Y \neq g(X)}] = \mathbb{P}_{\mathcal{D}}[Y \neq g(X)]$$

classification error

probability of making an error

# Bayes classifier

What is the **optimal performance**, regardless of the finiteness of the training data?

Def: The classifier  $g_* = \arg \min_g L_{\mathcal{D}}(g)$  is called the **Bayes classifier**

Claim:

$$g_*(x) = \arg \max_{y \in \{0,1\}} \mathbb{P}(Y = y | X = x)$$

# Proof of the Bayes classifier

**Claim 1:**  $\forall x \in \mathcal{X}, h_*(x) \in \arg \min_{y \in \mathcal{Y}} \mathbb{P}(Y \neq y | X = x) \implies h_* \in \arg \min_{h: \mathcal{X} \rightarrow \mathcal{Y}} L_{\mathcal{D}}(h)$

$$\begin{aligned} L_{\mathcal{D}}(h) &= \mathbb{E}_{X,Y}[1_{Y \neq h(X)}] = \mathbb{E}_X[\mathbb{E}_{Y|X}[1_{Y \neq h(X)} | X]] \\ &= \mathbb{E}_X[\mathbb{P}(Y \neq h(X) | X)] \\ &\geq \mathbb{E}_X[\min_{y \in \mathcal{Y}} \mathbb{P}(Y \neq y | X)] \\ &= \mathbb{E}_X[\mathbb{P}(Y \neq h_*(X) | X)] = \mathbb{E}_{X,Y}[1_{Y \neq h_*(X)}] = L_{\mathcal{D}}(h_*) \end{aligned}$$

**Claim 2:**  $g_*(x) = \arg \min_{y \in \mathcal{Y}} \mathbb{P}(Y \neq y | X = x)$

$$g_*(x) = \arg \max_{y \in \mathcal{Y}} \mathbb{P}(Y = y | X = x) = \arg \min_{y \in \mathcal{Y}} \mathbb{P}(Y \neq y | X = x)$$



# Bonus: a good regressor implies a good classifier

For all regression functions  $\eta : \mathcal{X} \rightarrow \mathbb{R}$  we can define a classifier as

$$\begin{aligned} \mathcal{X} &\rightarrow \{0,1\} \\ g_\eta : x &\mapsto 1_{\eta(x) \geq 1/2} \end{aligned}$$

**Claim:**

$$L_{\mathcal{D}}^{\text{classif}}(g_\eta) - L_{\mathcal{D}}^{\text{classif}}(g^*) \leq 2\sqrt{L_{\mathcal{D}}^{\ell_2}(\eta) - L_{\mathcal{D}}^{\ell_2}(\eta^*)}$$

Where  $L_{\mathcal{D}}^{\text{classif}}(g_\eta) = \mathbb{E}_{\mathcal{D}}[1_{g(X) \neq Y}]$ ,  $L_{\mathcal{D}}^{\ell_2}(f) = \mathbb{E}_{\mathcal{D}}[(Y - f(X))^2]$  and  $\eta_* = \arg \min_{\eta} L_{\mathcal{D}}^{\ell_2}(\eta)$

➡ If  $\eta$  is good for regression then  $g_\eta$  is good for classification too (converse is not true)