

# Neural Networks: Basic Structure, Representation Power

Machine Learning Course - CS-433

Nov 9, 2021

Nicolas Flammarion





# A lot of hype for NNs: perform very well in practice but are still theoretically misunderstood





# Neural Networks: motivation

Supervised learning : we observe some data  $S_{\text{train}} = \{x_i, y_i\}_{i=1}^n \in \mathcal{X} \times \mathcal{Y}$


➡ given a new  $x$ , we want to predict its label  $y$

Linear predictions: it works well only when used with **good features**

Data representation:


- Before: domain experts derived efficient features for a particular application
- Now: neural networks directly learn the features from the data

# NNs: the basic structure



- $x_i^{(l)}$ : value for the  $i$ -th node of the  $l$ -th layer
- Each layer has the same number of nodes
- Feedforward network - there is no feedback loop
- Same NN for regression and classification - only the output layer changes

# Fully connected NNs: each node is connected to all the nodes in the previous layer



Parameters of the network to be learnt:

- $w_{i,j}^{(l)}$ : weight associated with the edge going from node  $i$  in layer  $l - 1$  to node  $j$  in layer  $l$
- $b_j^{(l)}$ : bias term associated with node  $j$  in layer  $l$

# The function value at the $l^{th}$ layer is defined recursively


The output is given by

$$x_j^{(l)} = \phi \left( \sum_{i=1}^K x_i^{(l-1)} w_{i,j}^{(l)} + b_j^{(l)} \right)$$

$\phi$ : activation function:

- Sigmoid  $\frac{1}{1 + e^{-x}}$
- RELU  $\max\{0, x\}$

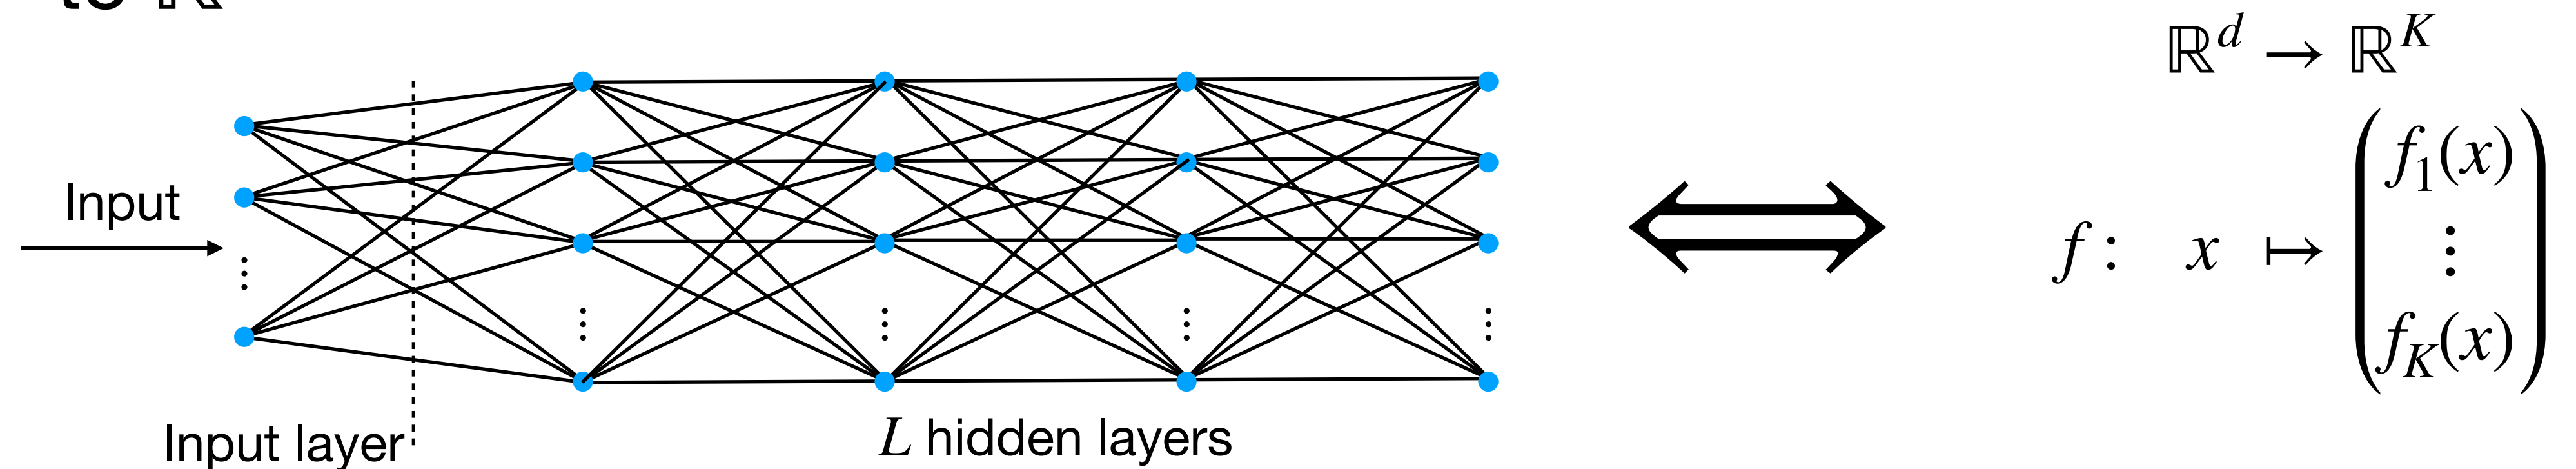
Key fact:  $\phi$  is non-linear - otherwise NNs only represent linear functions



Sigmoid function

# The NN transforms the input into a more suitable representation

The NN can be decomposed in two. The first part represents a function from  $\mathbb{R}^d$  to  $\mathbb{R}^K$




This function is defined by

- The biases  $\{b_i^{(l)}\}_{i \in [K], l \in [L]}$  and weights  $\{w_{i,j}^{(l)}\}_{i,j \in [K], l \in [L]}$  we learn  
     $\Rightarrow O(K^2L)$  parameters
- The activation function  $\phi$  we pick

In practice: both  $L$  and  $K$  are large - overparametrized NNs

# The last layer performs the desired ML task



A suitable representation of the data in hands, the last layer only performs a linear regression or classification step



# The three main challenges of deep learning

- **Expressive power** of NNs: why are the function we are interested in so **well approximated** by NNs?
- **Success of naive optimisation**: why does **gradient descent** lead to a good local minimum?
- **Generalization miracle**: why is there **no overfitting** with so many parameters?

# Barron's Approximation result

Let  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  and define  $\hat{f}(\omega) = \int_{\mathbb{R}^d} f(x) e^{-i\omega^\top x} dx$  its Fourier transform

Assumption:  $\int_{\mathbb{R}^d} |\omega| |\hat{f}(\omega)| d\omega \leq C$  (smoothness assumption)

Claim: For all  $n \geq 1$ , it exists a function  $f_n$  of the form

$$f_n(x) = \sum_{j=1}^n c_j \phi(x^\top w_j + b_j) + c_0$$

so that

$$\int_{|x| \leq r} (f(x) - f_n(x))^2 dx \leq \frac{(2Cr)^2}{n}$$

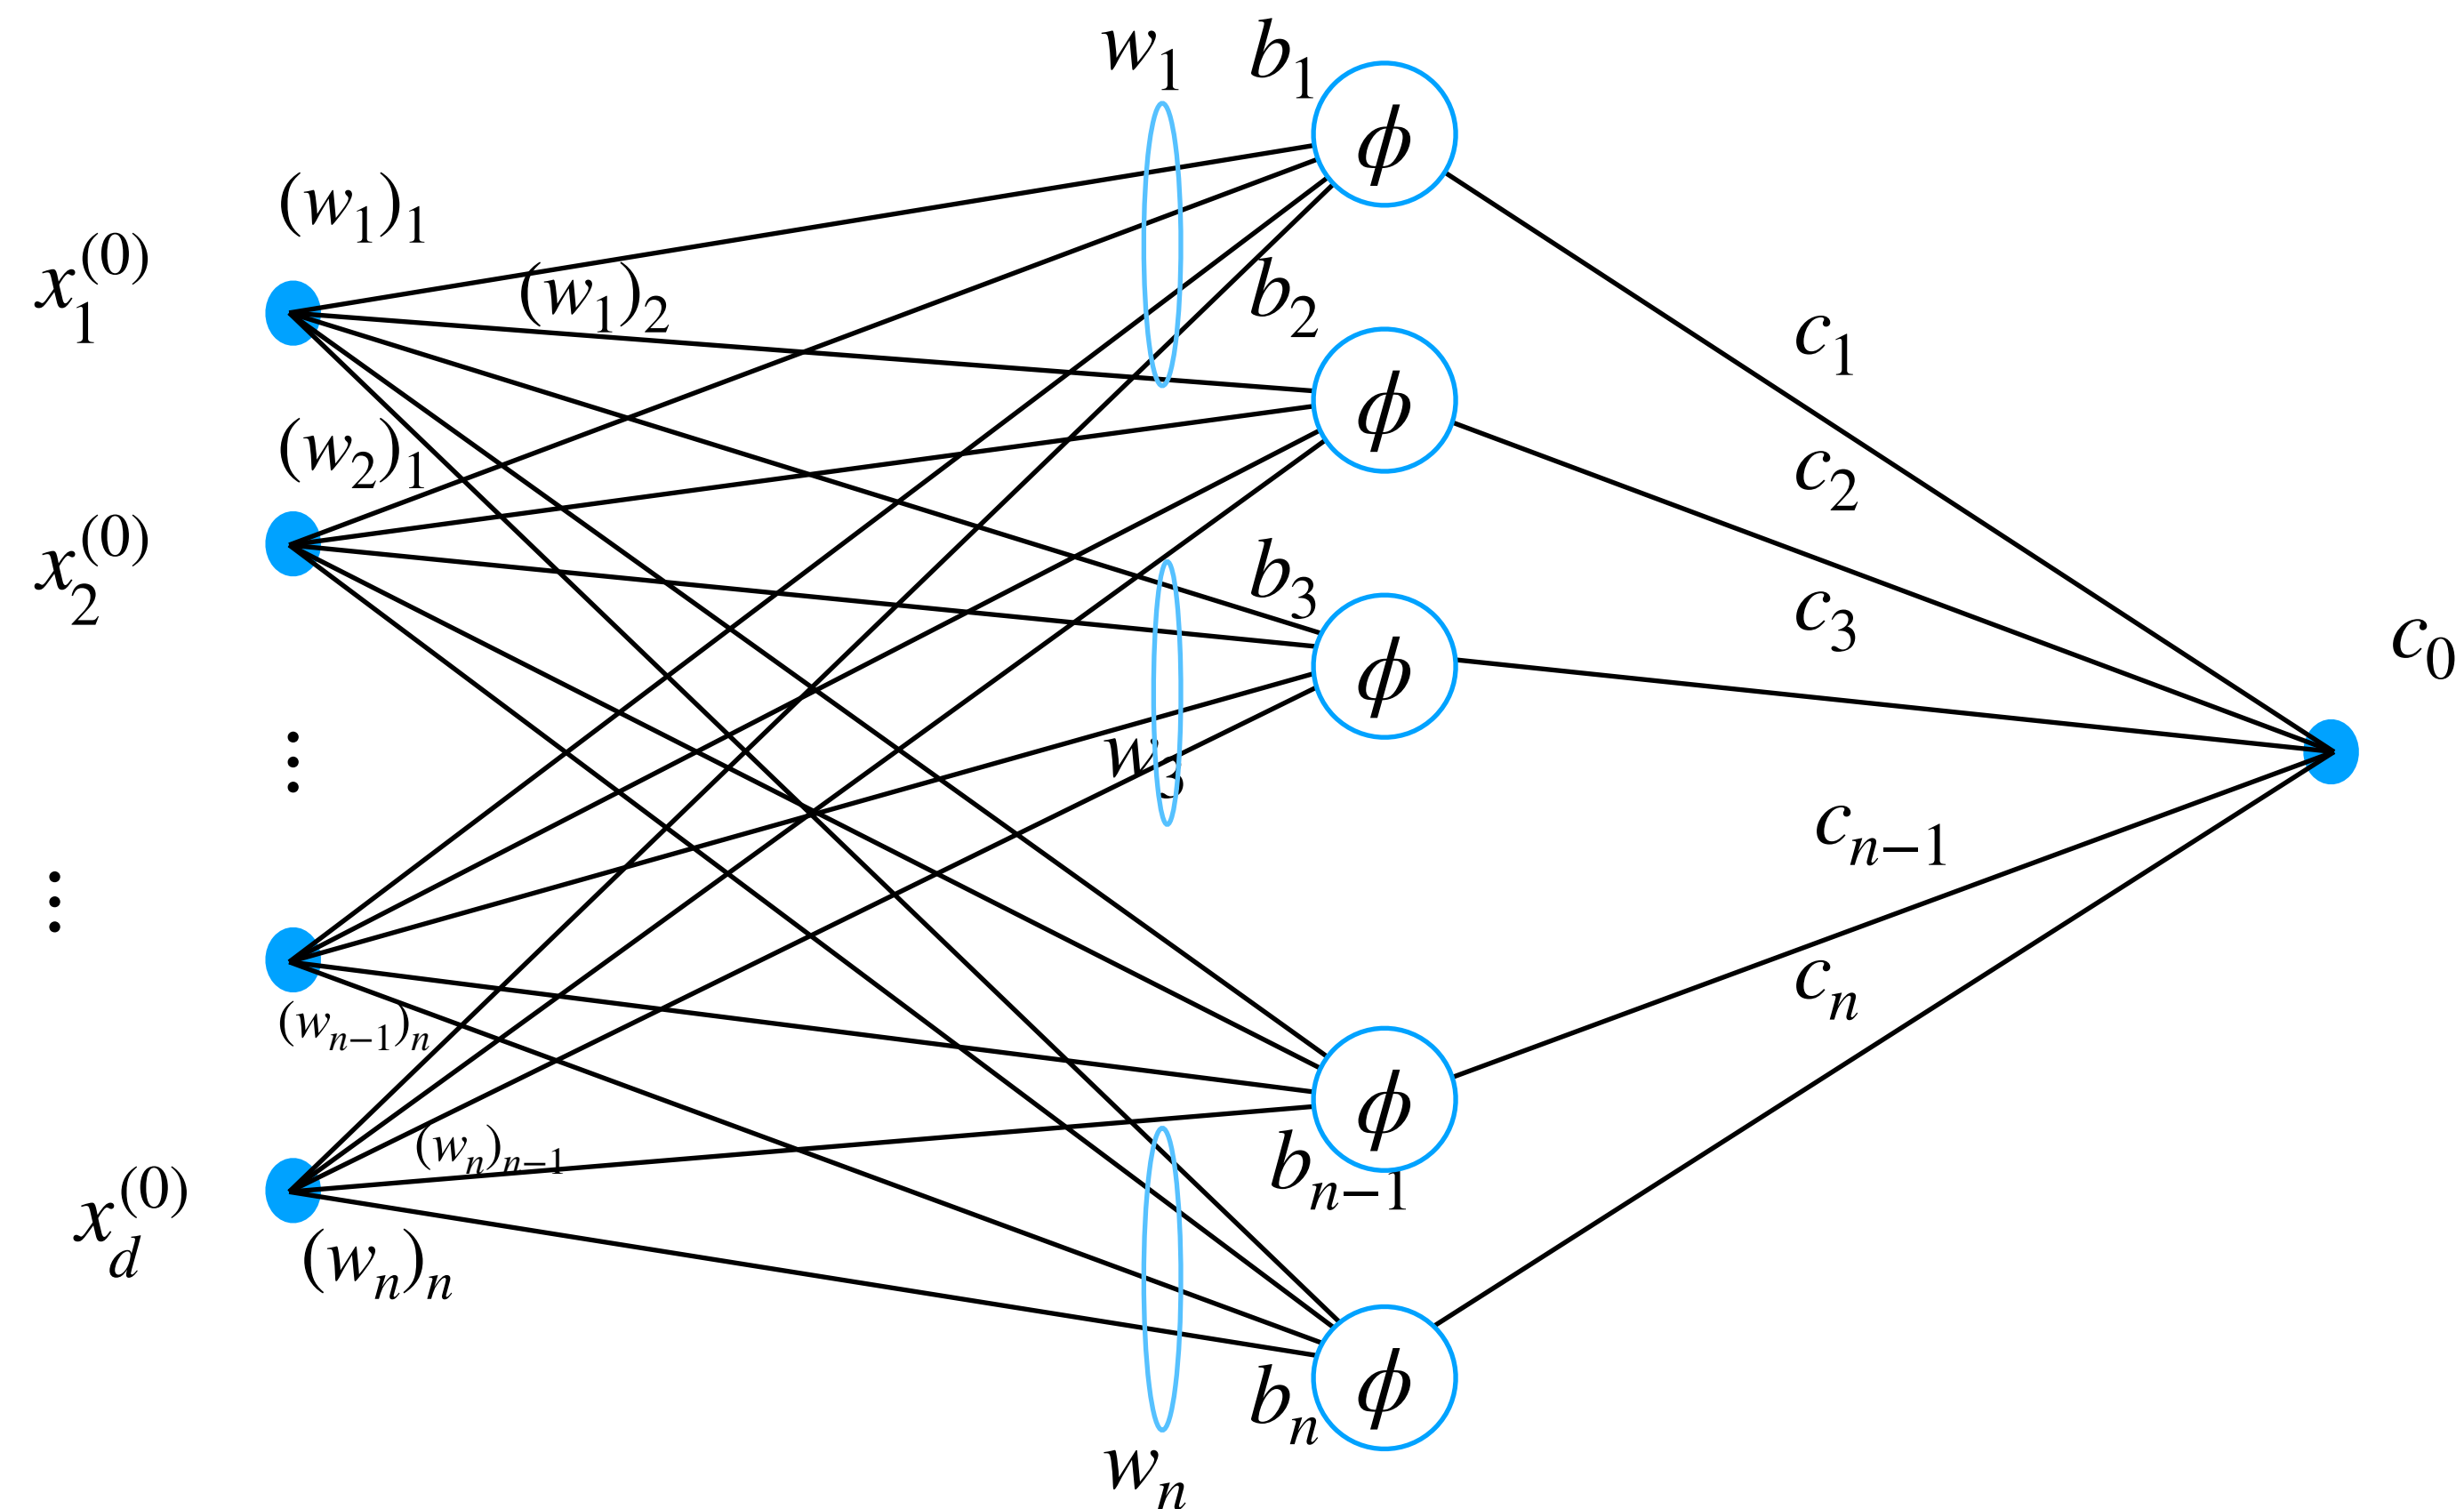
# All sufficiently smooth function can be approximated by a one-hidden-layer NN

$$\int_{|x| \leq r} \left( f(x) - f_n(x) \right)^2 dx \leq \frac{(2Cr)^2}{n}$$

- The more neurons we allow, the smaller the error
- The smoother the function is (the smaller  $C$ ), the smaller the error
- The larger the domain (the larger  $r$ ), the worse the error
- Approximation in average (in  $\ell_2$ -norm)
- For any “sigmoid-like” activation function

The function  $f_n$  is a one-hidden-layer NN with  $n$  nodes

$$f_n(x) = \sum_{j=1}^n c_j \phi(x^\top w_j + b_j) + c_0$$






# Proof by picture

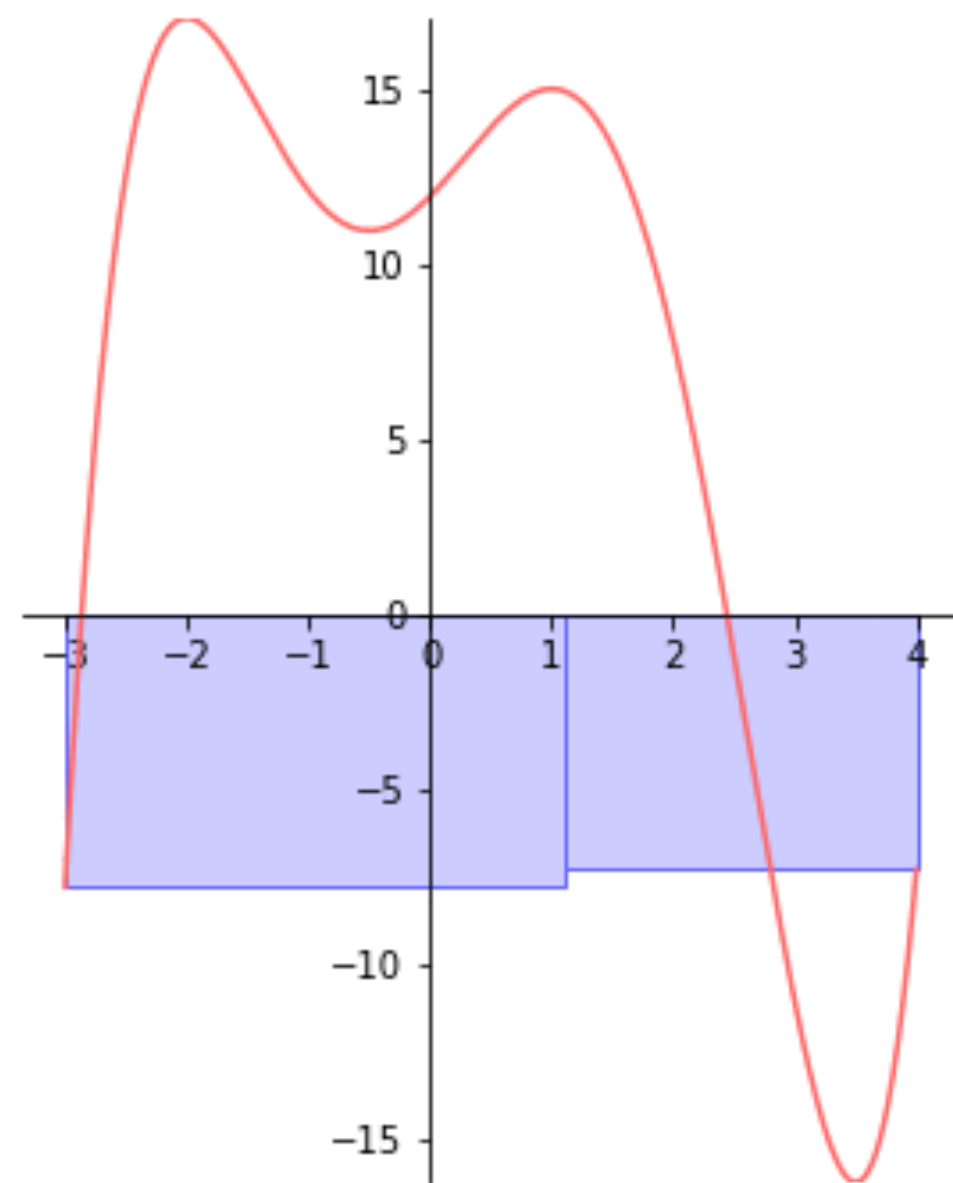
Simple and intuitive explanation of a slightly different result:

“A NN with sigmoid activation and at most two hidden layers can approximate well a smooth function in average, i.e, in  $\ell_1$ -norm”


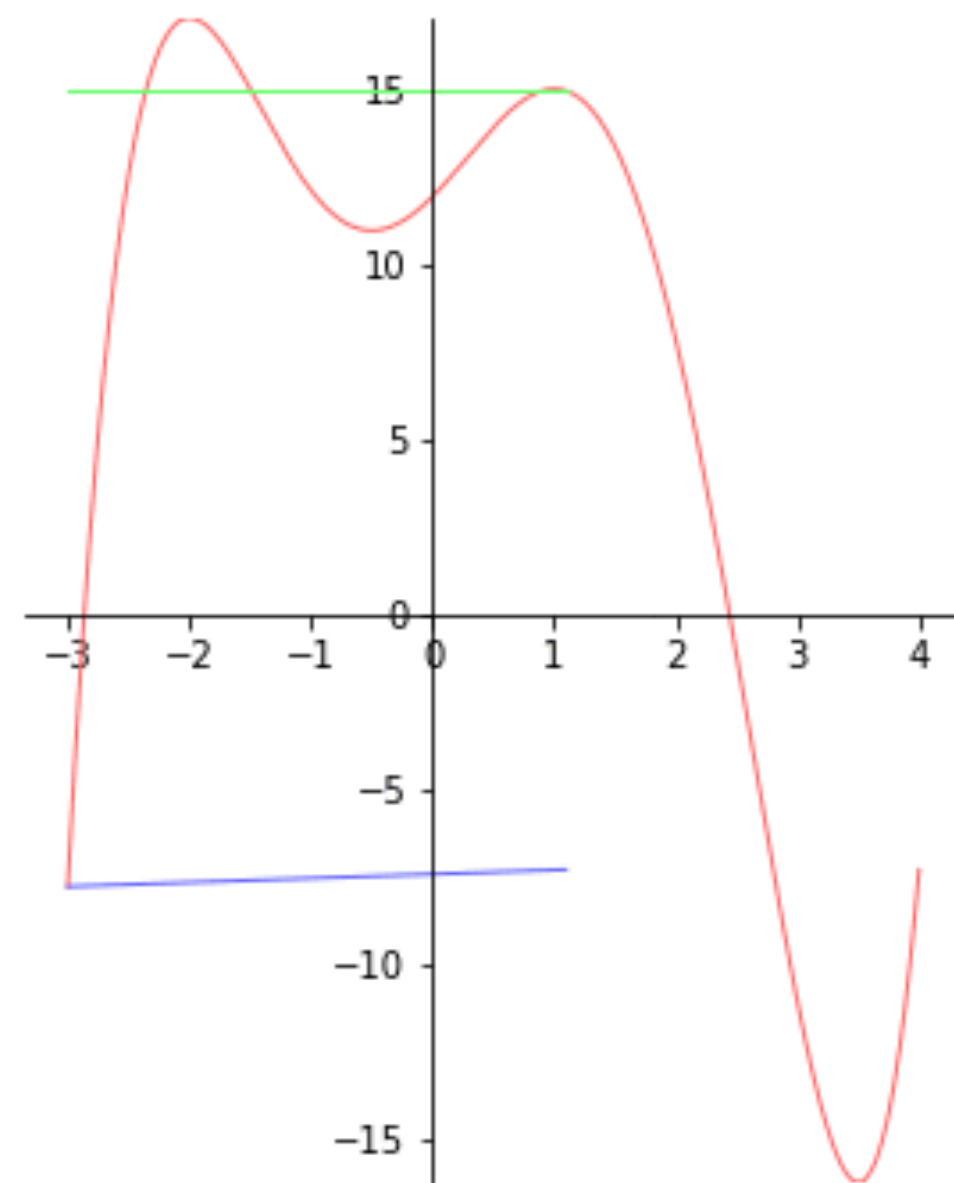
Consider a function  $f : \mathbb{R} \rightarrow \mathbb{R}$  on a bounded domain



# Approximation of the function by a sum of rectangles



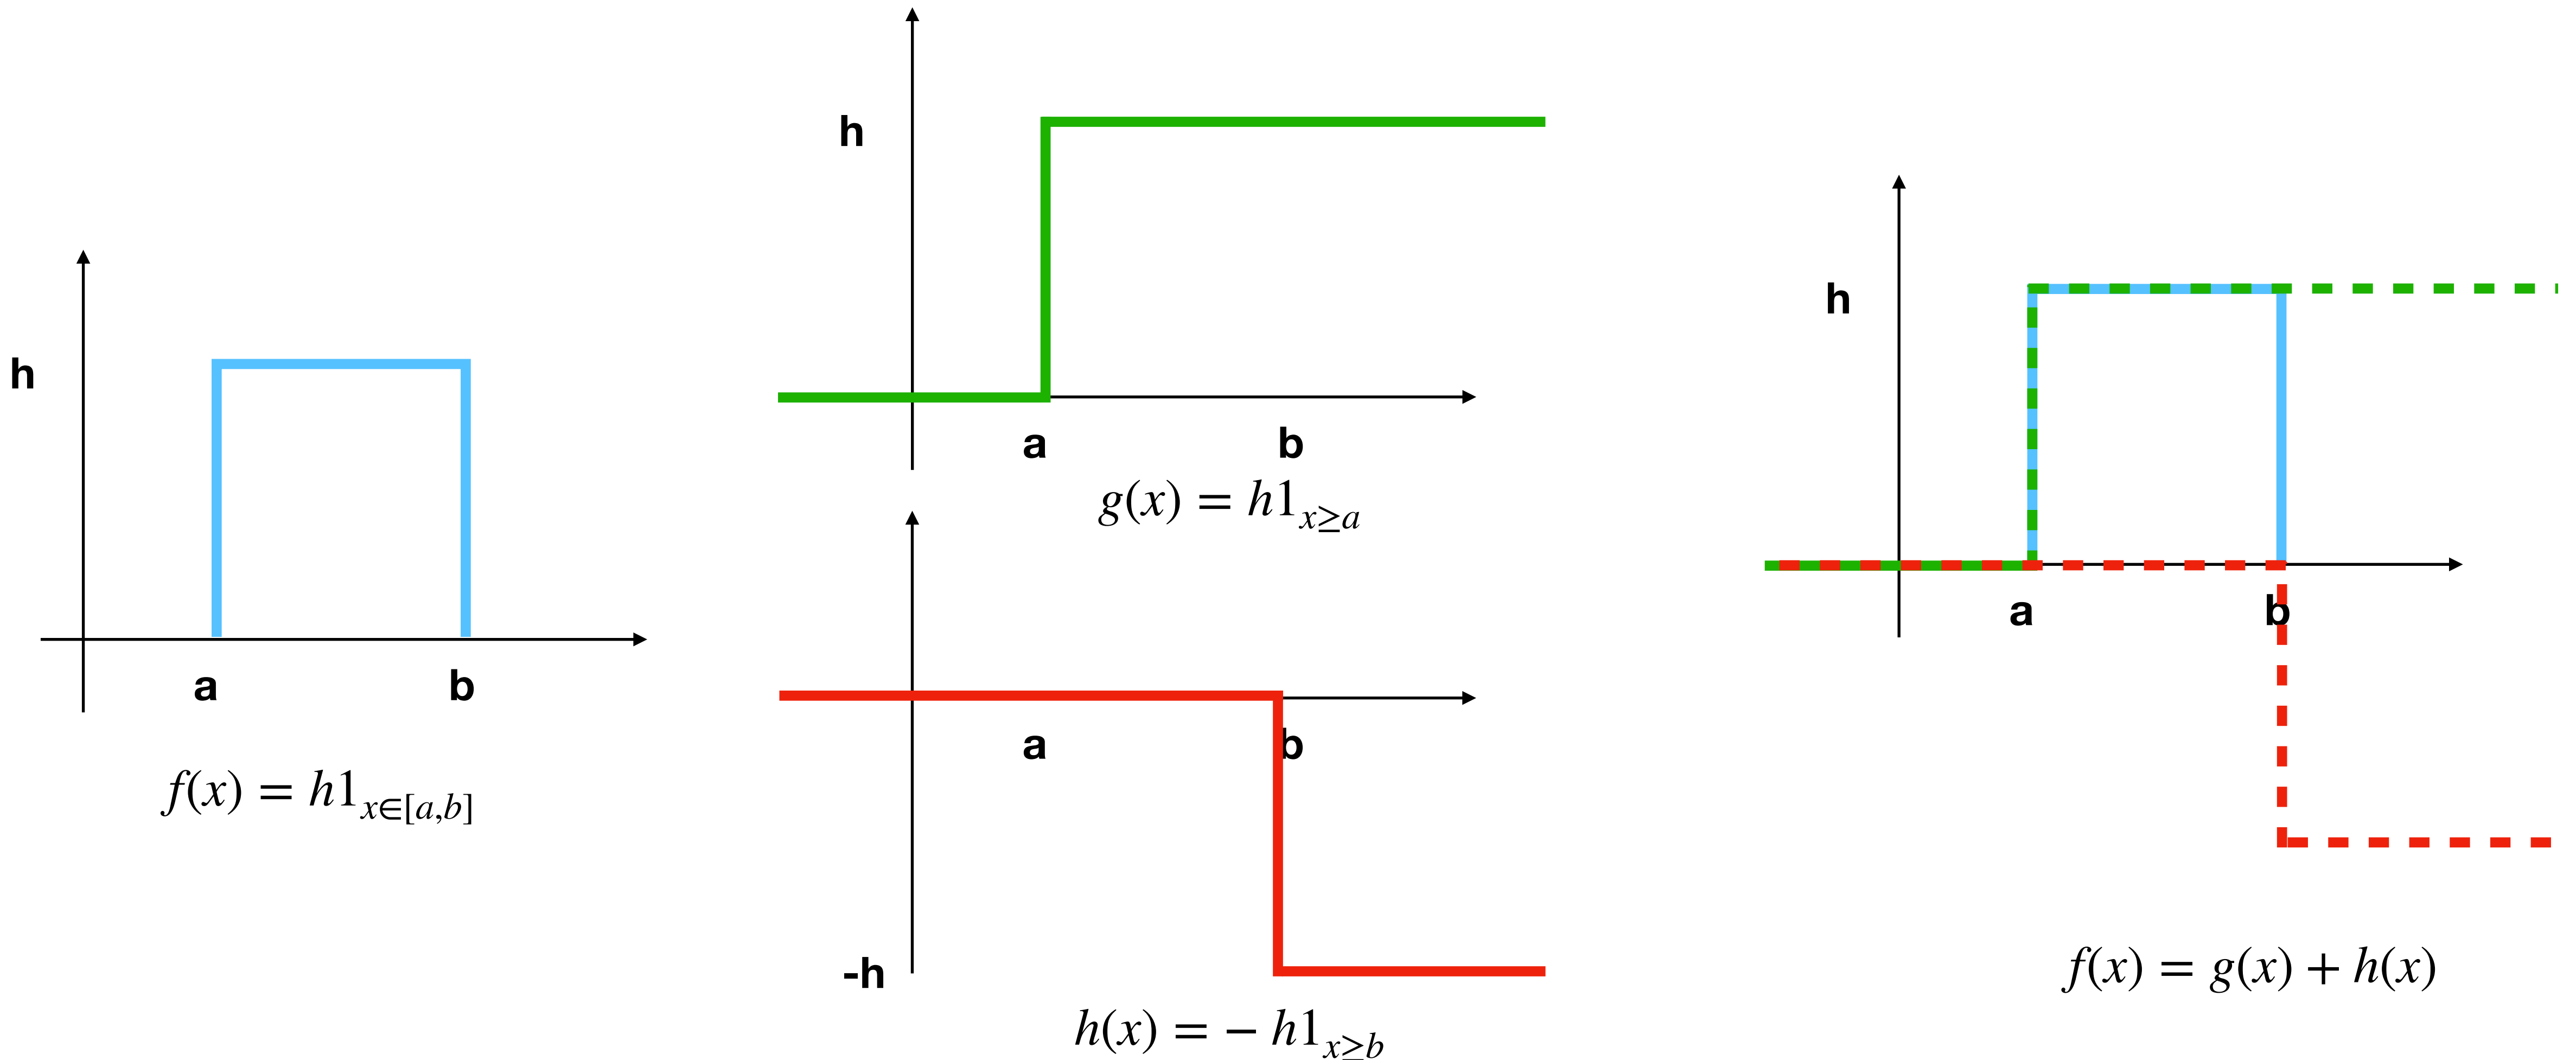
From below



From above

The function is Riemann integrable - can be approximated arbitrarily closely by “lower” and “upper” sums of rectangle

# A rectangle is equal to the sum of two step functions



# Approximate a step function with a sigmoid

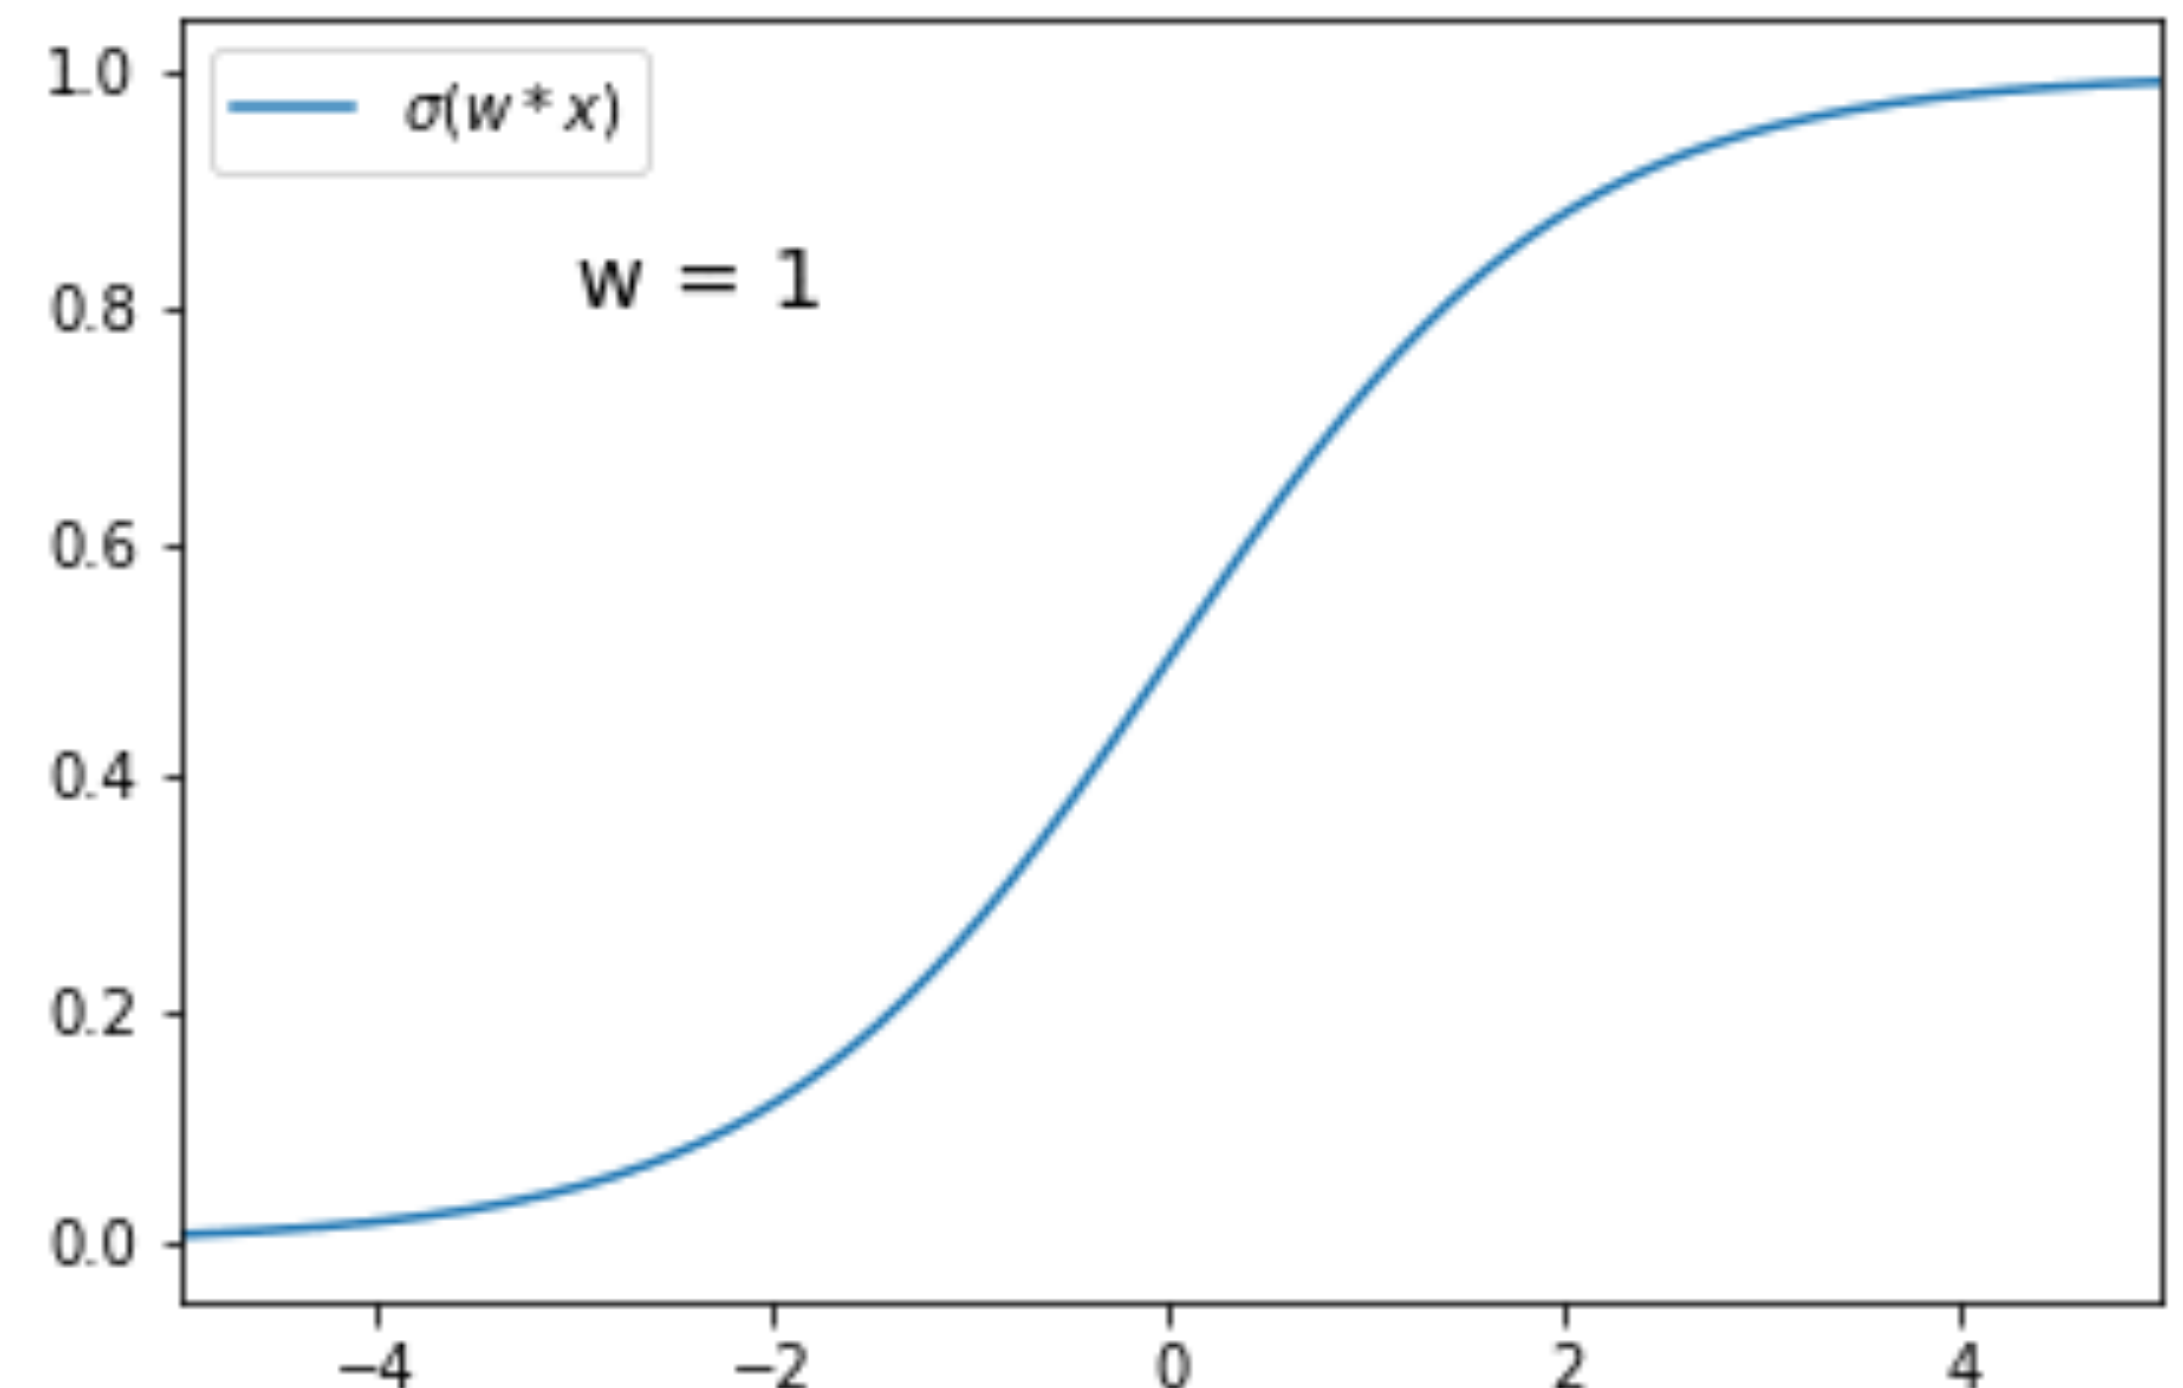
$$\tilde{\phi}(x) = \phi(w(x - b))$$

By setting:

- $b$ : where the transition happens
- $w$ : makes the transition steeper

Derivative:  $\tilde{\phi}'(b) = w/4$




➡ The width of the transition is  $O(4/w)$





# Approximation of the rectangle

$$h(\phi(w(x - a)) - \phi(w(x - b)))$$




# Conclusion in the 1D case

1. Approximate the function in the Riemann sense by a sum of  $k$  rectangles
2. Approximate each rectangle, by means of two nodes in the hidden layer of a nn
3. Compute the sum (with appropriate sign) of all the hidden layers at the output node
  - ➡ NN with one hidden layer containing  $2k$  nodes for a Riemann sum with  $k$  rectangles

Rmk:

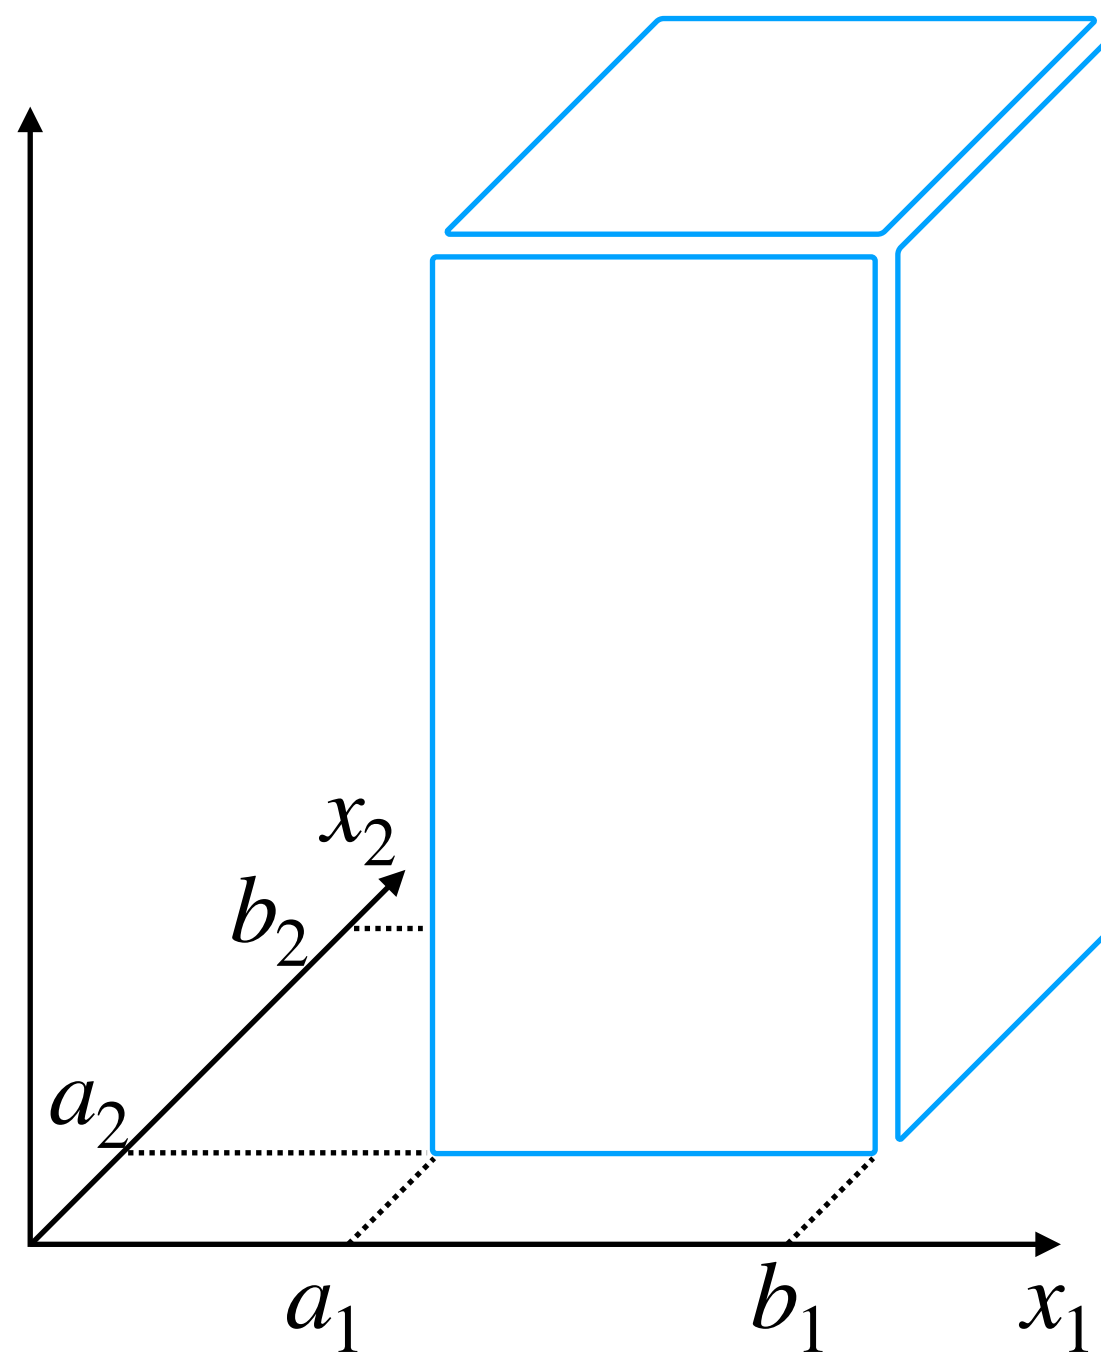
- Same intuition hold for any sigmoid like function
- Only intuition, not quantitative
- Need the weights  $w$  to be large



# Larger dimension: $d = 2$

Same idea:

- Approximate the function by 2D rectangles
- Approximate a 2D rectangle by sigmoids





# Two sigmoids can approximate an infinite rectangle

$$(x_1, x_2) \mapsto \phi(w(x_1 - a_1)) - \phi(w(x_1 - b_1))$$

Rectangle:

- going from  $a_1$  to  $b_1$  in the  $x_1$  direction
- unbounded in the  $x_2$  direction



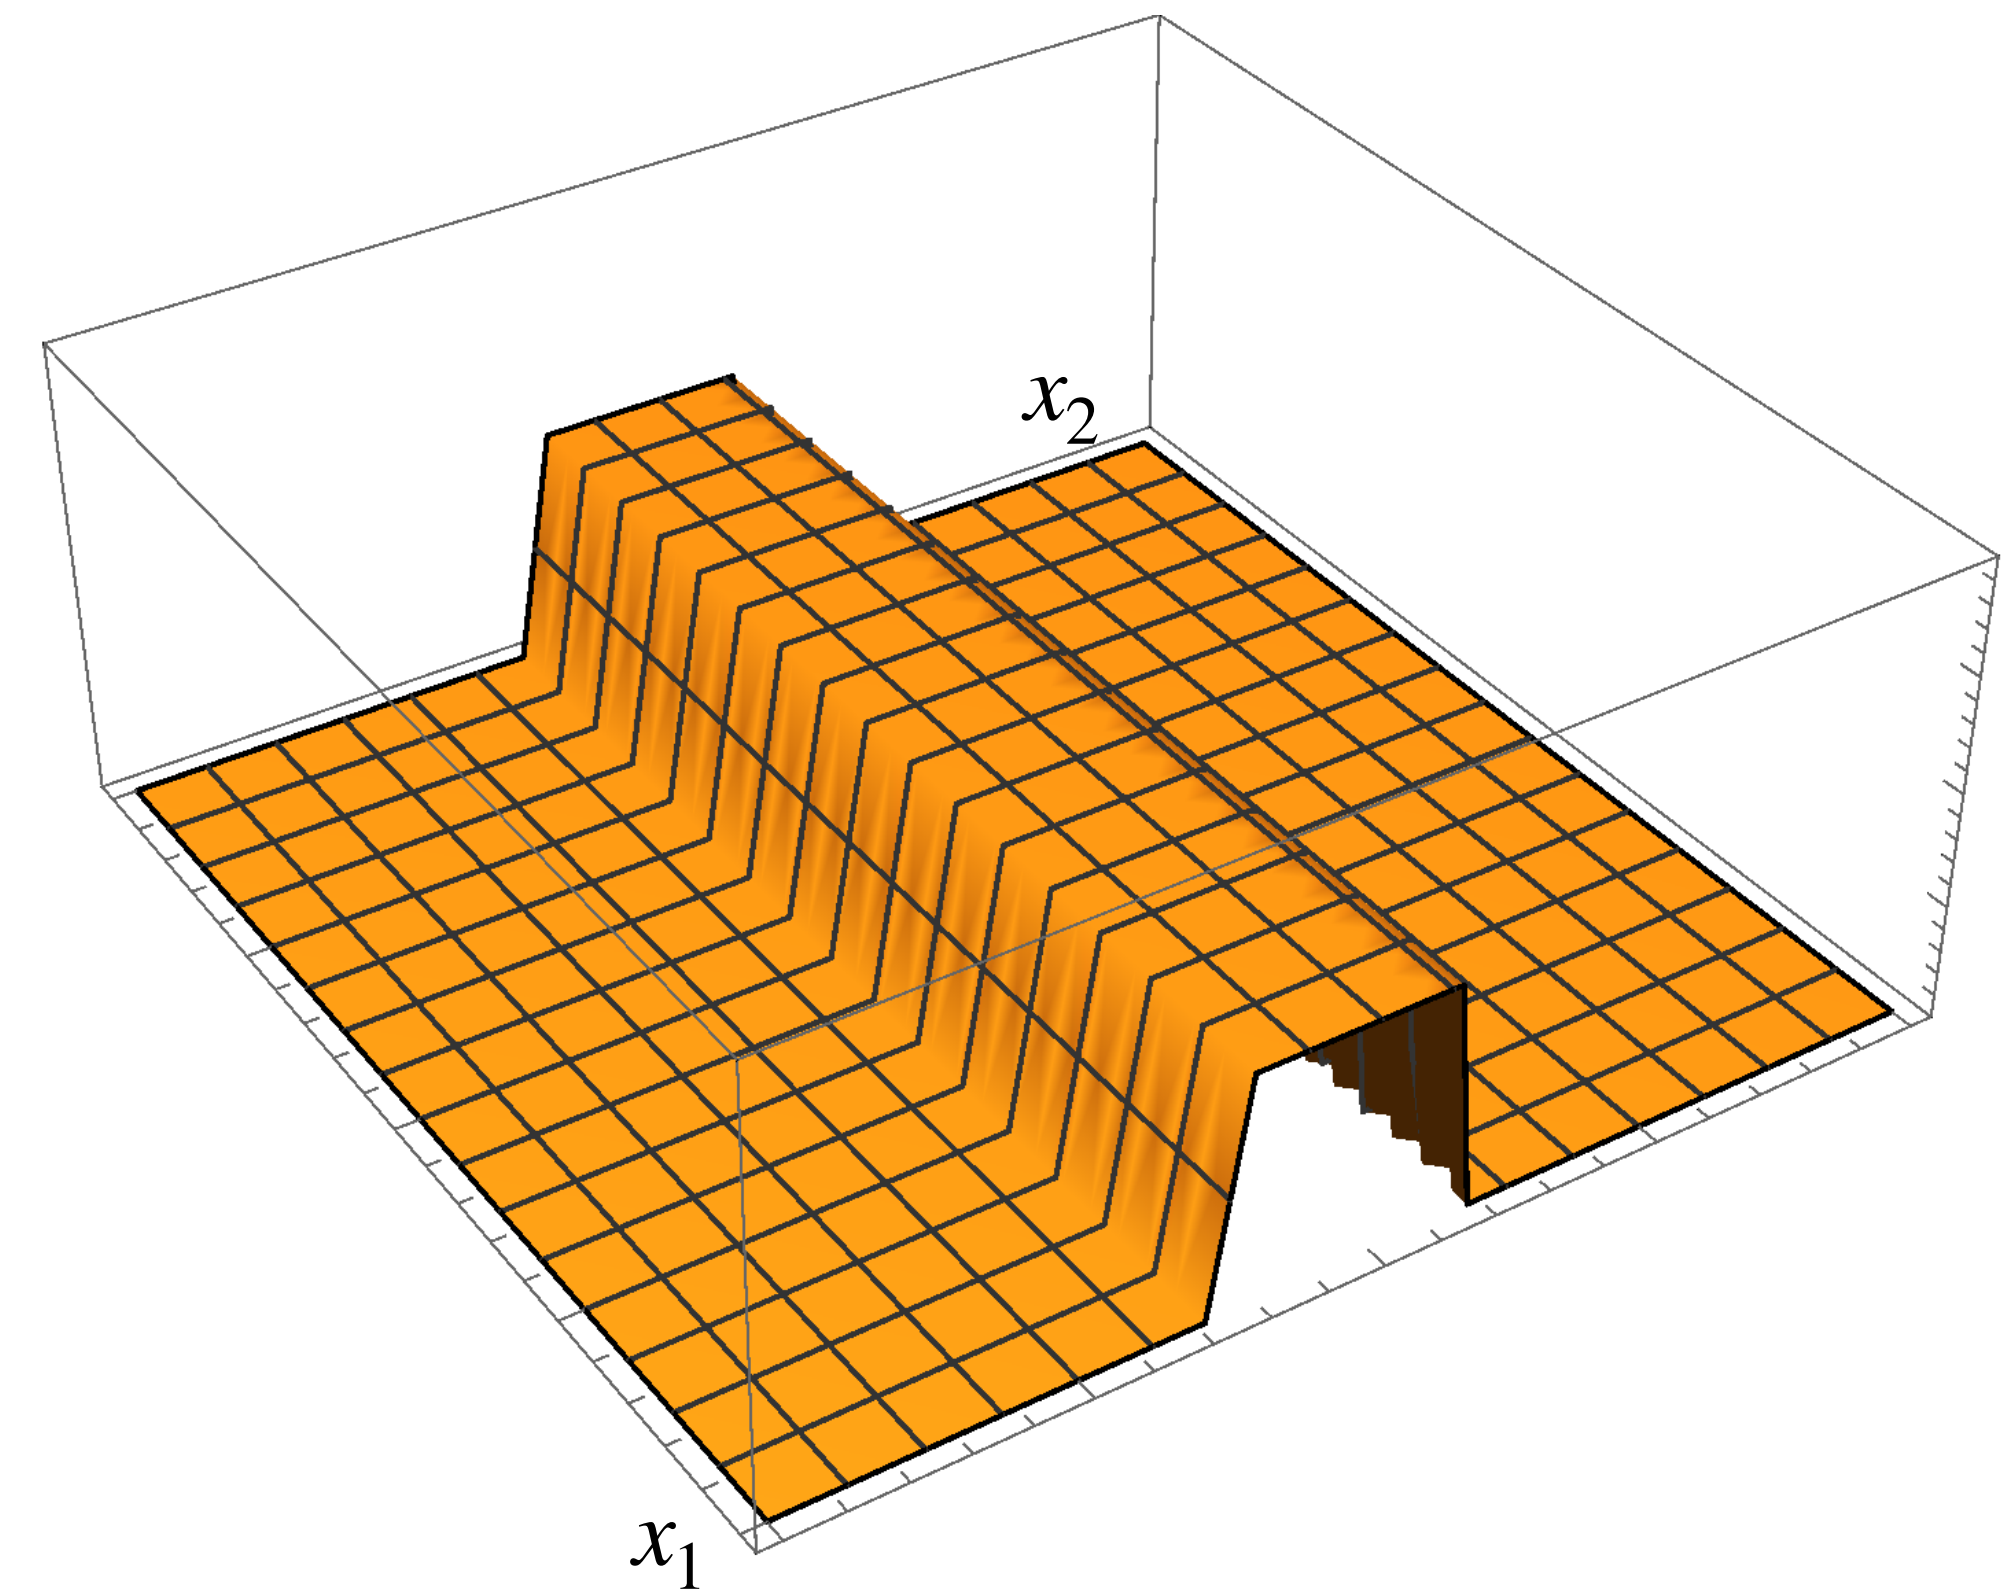



# Two sigmoids can approximate an infinite rectangle

$$(x_1, x_2) \mapsto \phi(w(x_2 - a_2)) - \phi(w(x_2 - b_2))$$

Rectangle:

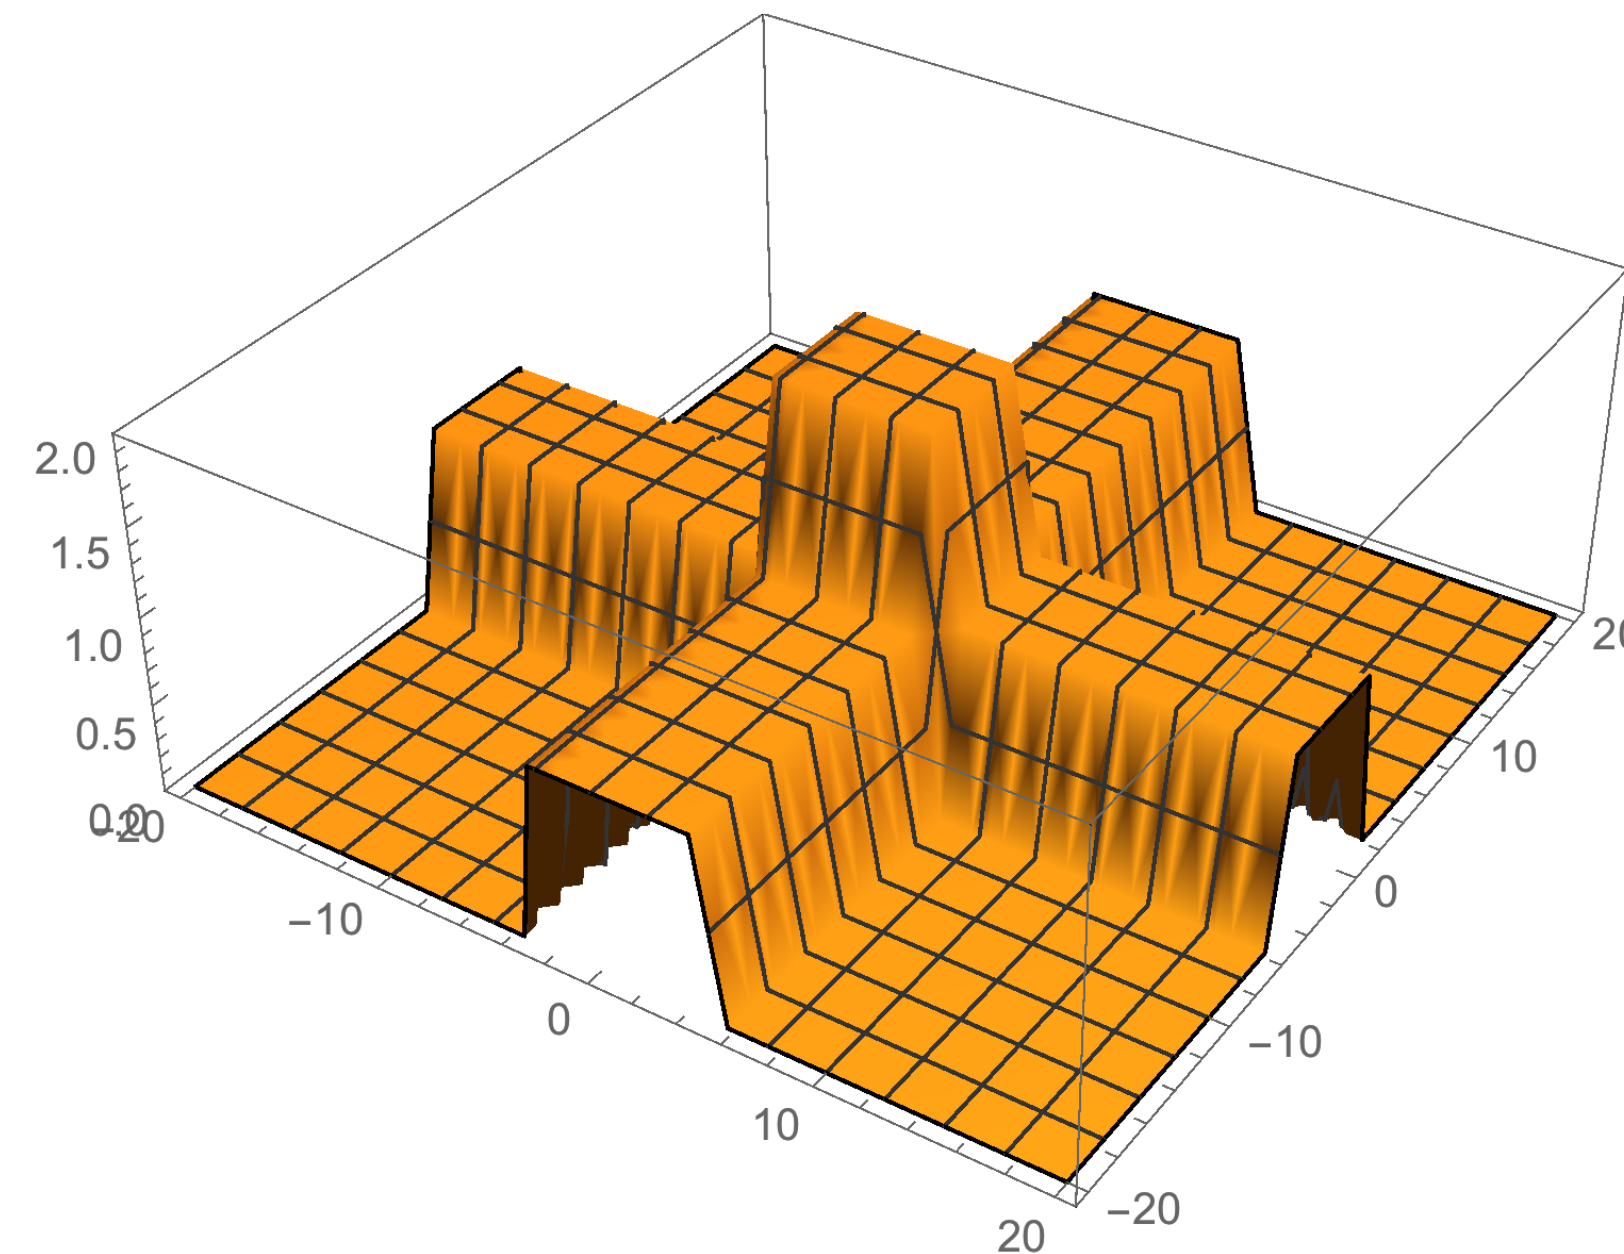

- going from  $a_2$  to  $b_2$  in the  $x_2$  direction
- unbounded in the  $x_1$  direction



# Four sigmoids can approximate a cross

$$(x_1, x_2) \mapsto \phi(w(x_1 - a_1)) - \phi(w(x_1 - b_1)) + \phi(w(x_2 - a_2)) - \phi(w(x_2 - b_2))$$

➡ close to what we want with the exception of the two infinite “arms”





How to get rid of the crossed arms?

# The sigmoid can threshold the unwanted infinite arms

Threshold the function value would remove the arms

It is equivalent to compose it with  $1_{y \geq c}$  for  $c \in (1,2]$

- ➡ approximate  $1_{y \geq c}$  by a sigmoid with large weight  $w$  and appropriate bias (e.g.,  $3w/2$ )





# Point-wise approximations

Let  $f$  be a continuous function on  $[0,1]$

Different  $\ell_\infty$ -approximation results exist:

- Polynomials: Stone Weierstrass theorem:

$\forall \varepsilon > 0, \exists p \in \mathbb{R}[X],$


$$\sup_{x \in [0,1]} |f(x) - p(x)| \leq \varepsilon$$

- Piecewise linear function (Shektman, 1982)

$$q(x) = \sum_{i=1}^m (a_i x + b_i) 1_{r_{i-1} \leq x < r_i}$$

with  $a_i r_i + b_i = a_{i+1} r_i + b_{i+1}$

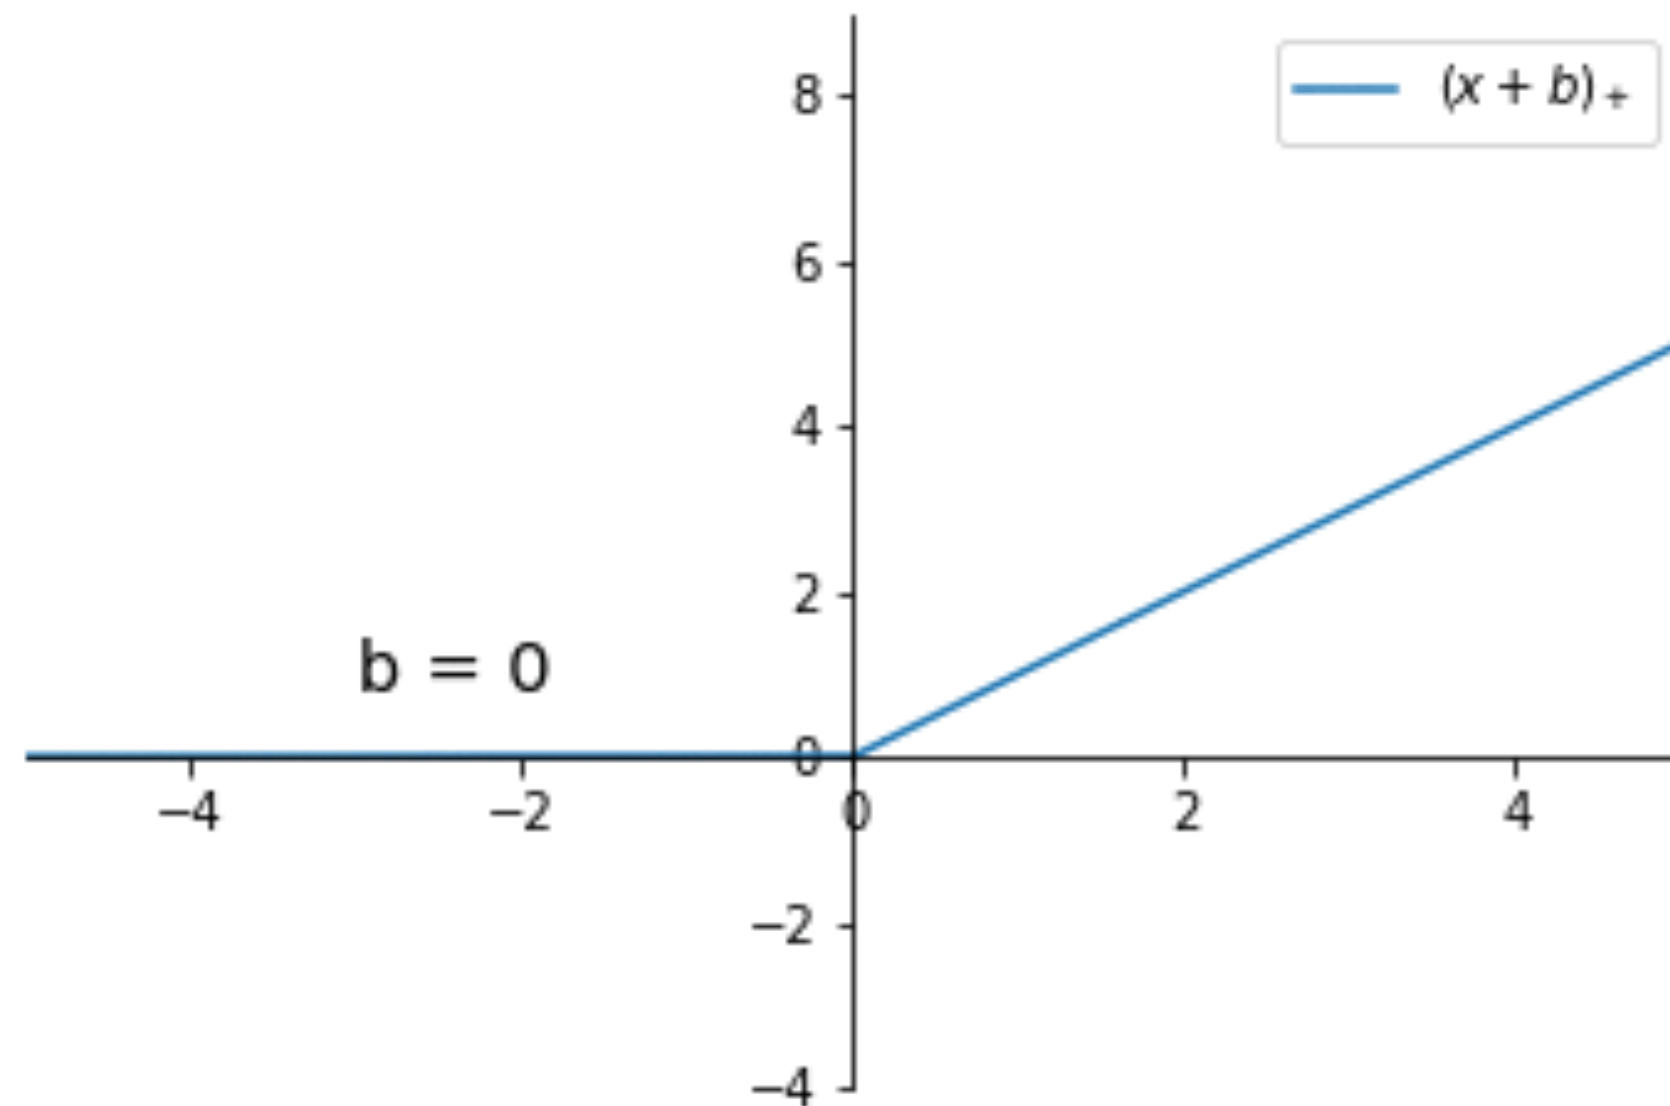
➡ How to approximate such functions with a NN?



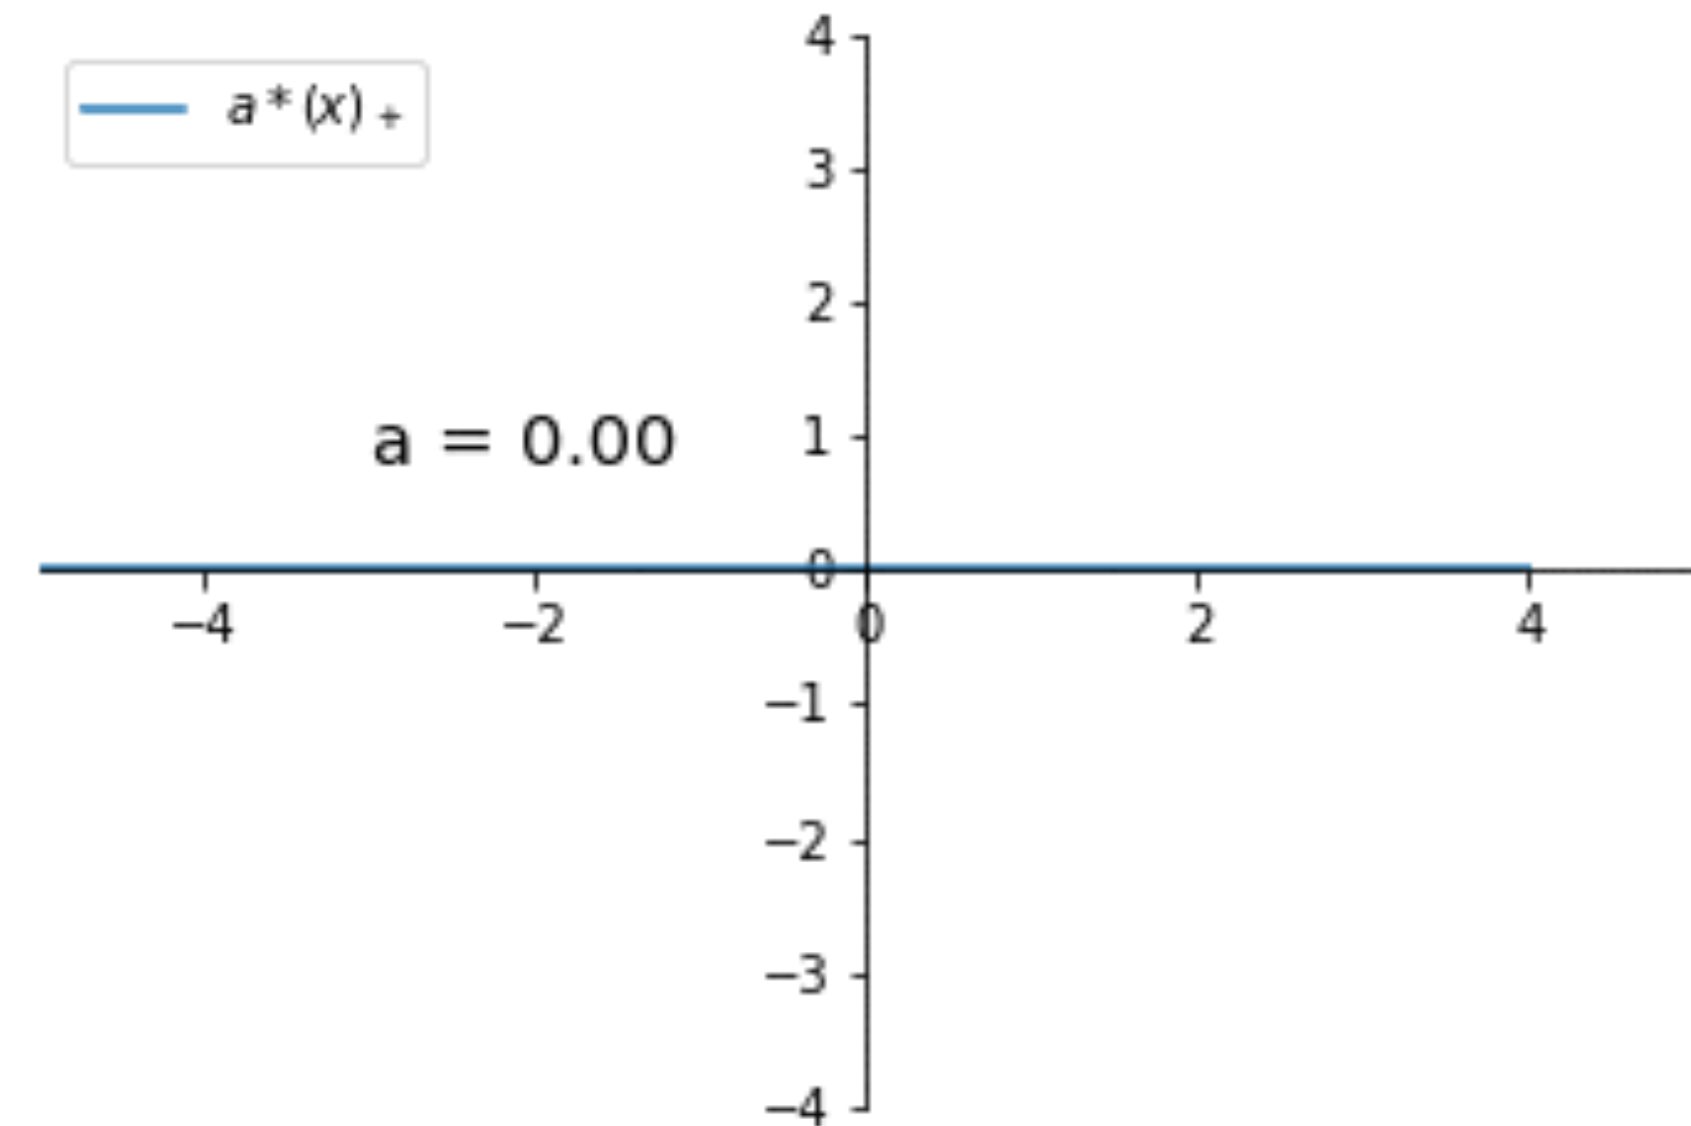


# Rectified linear unit - RELU

$$(x)_+ = \max\{0, x\}$$



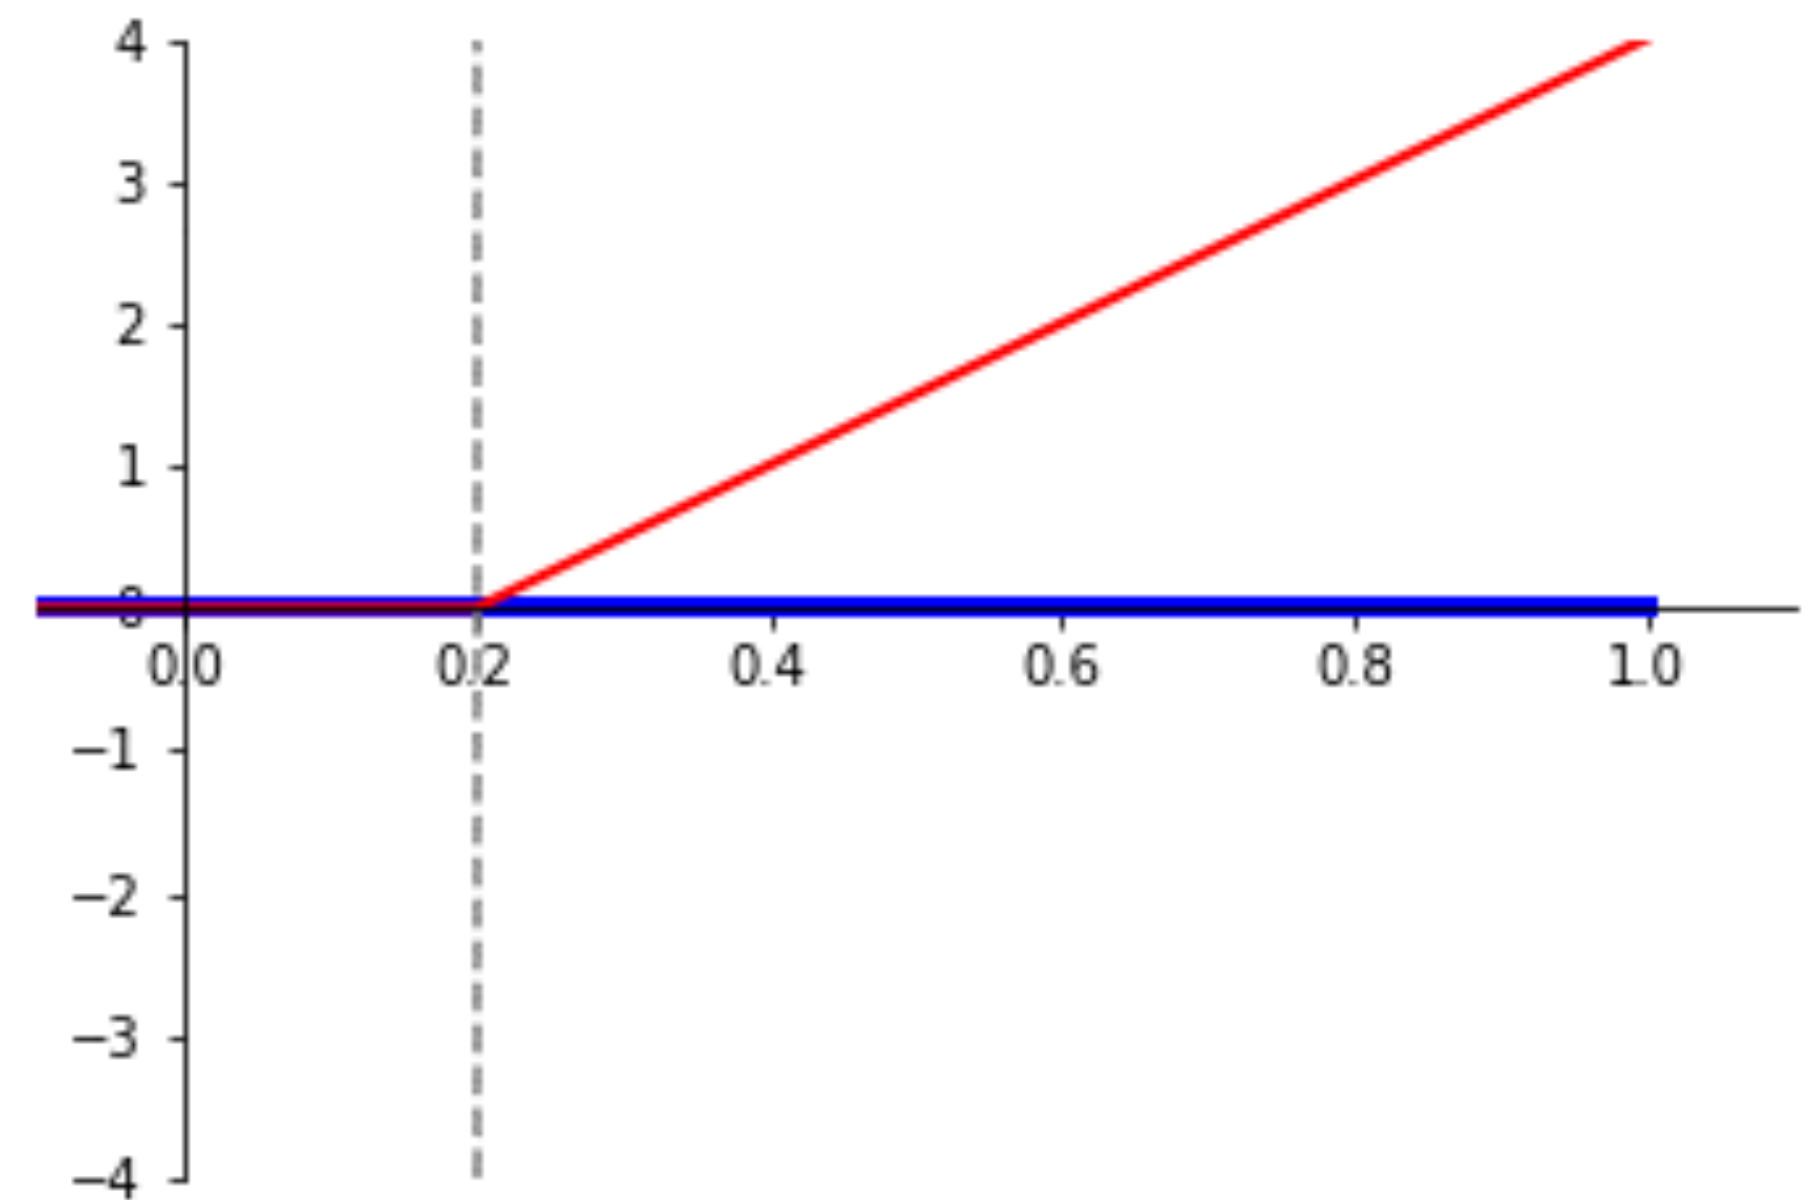
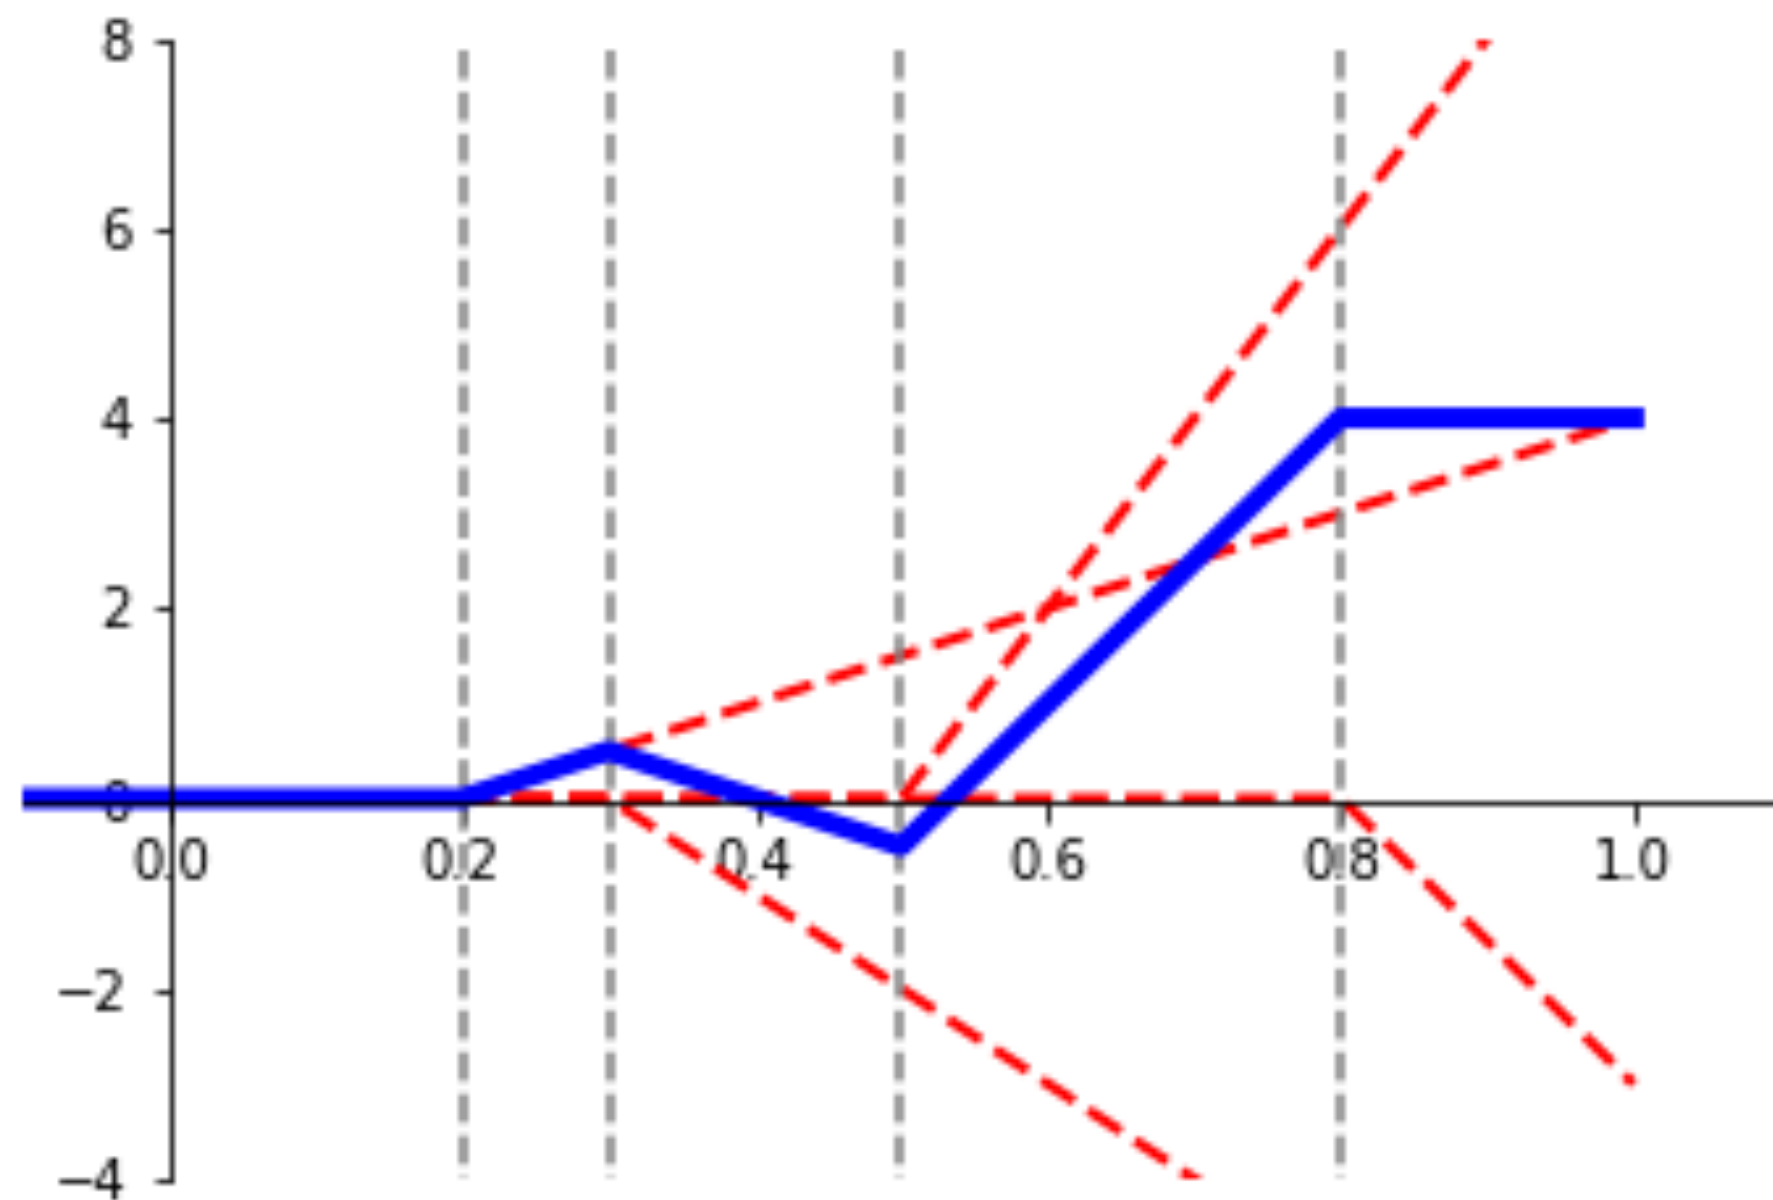
The bias  $b$  determines where the kink is



The weight  $a$  determines the slope

# Linear combinations of RELUs are piecewise linear functions

$\sum_{i=1}^m \tilde{a}_i (x - \tilde{b}_i)_+$  is a piecewise linear function



# Piecewise linear functions can be written as combination of RELU

Claim 1:  $q$  can be rewritten as

$$q(x) = \tilde{a}_1 x + \tilde{b}_1 + \sum_{i=2}^m \tilde{a}_i (x - \tilde{b}_i)_+$$


where  $\tilde{a}_1 = a_1$ ,  $\tilde{b}_1 = b_1$ ,  $a_i = \sum_{j=1}^i \tilde{a}_j$  and  $\tilde{b}_i = r_{i-1}$

Claim 2:  $q$  can be implemented as a one-hidden-layer NN with RELU activation. Each term corresponds to one node:

- Bias  $-\tilde{b}_i$
- Output weight  $\tilde{a}_i$

The term  $\tilde{a}_1 x + \tilde{b}_1$  also corresponds to one node:

- Bias  $\tilde{b}_1$ : bias of the output node
- Term  $\tilde{a}_1 x = \tilde{a}_1 (x)_+$  since  $x \in [0,1]$



# Proof of the equivalent formulation

$$q(x) = \sum_{i=1}^m (a_i x + b_i) 1_{r_{i-1} \leq x < r_i} \quad r(x) = \tilde{a}_1 x + \tilde{b}_1 + \sum_{i=2}^m \tilde{a}_i (x - \tilde{b}_i)_+$$

$$\tilde{a}_1 = a_1, \tilde{b}_1 = b_1 \text{ and } a_i = \sum_{j=1}^i \tilde{a}_j \text{ and } \tilde{b}_i = r_{i-1}$$

- For  $x \in [0, r_1]$

$$(\tilde{a}_1, \tilde{b}_1) = (a_1, b_1) \implies q(x) = a_1 x + b_1 = \tilde{a}_1 x + \tilde{b}_1 = r(x) \text{ because } \tilde{b}_2 = r_1$$

- For  $x \in [r_1, r_2]$ ,  $r(x) = \tilde{a}_1 x + \tilde{b}_1 + (a_2 - a_1)(x - r_1)_+$   
$$= a_1 x + b_1 + (a_2 - a_1)(x - r_1) = a_2 x + b_1 - (a_2 - a_1)r_1$$

$$r'(x) = a_2 \text{ and } r(r_1) = q(r_1) \text{ as shown above}$$

$$\implies r(x) = q(x) \text{ for } x \in [r_1, r_2]$$

# Proof by induction

Let's assume that  $r(x) = q(x)$  for  $x \in [0, r_{i-1}]$

For  $x \in [r_{i-1}, r_i]$

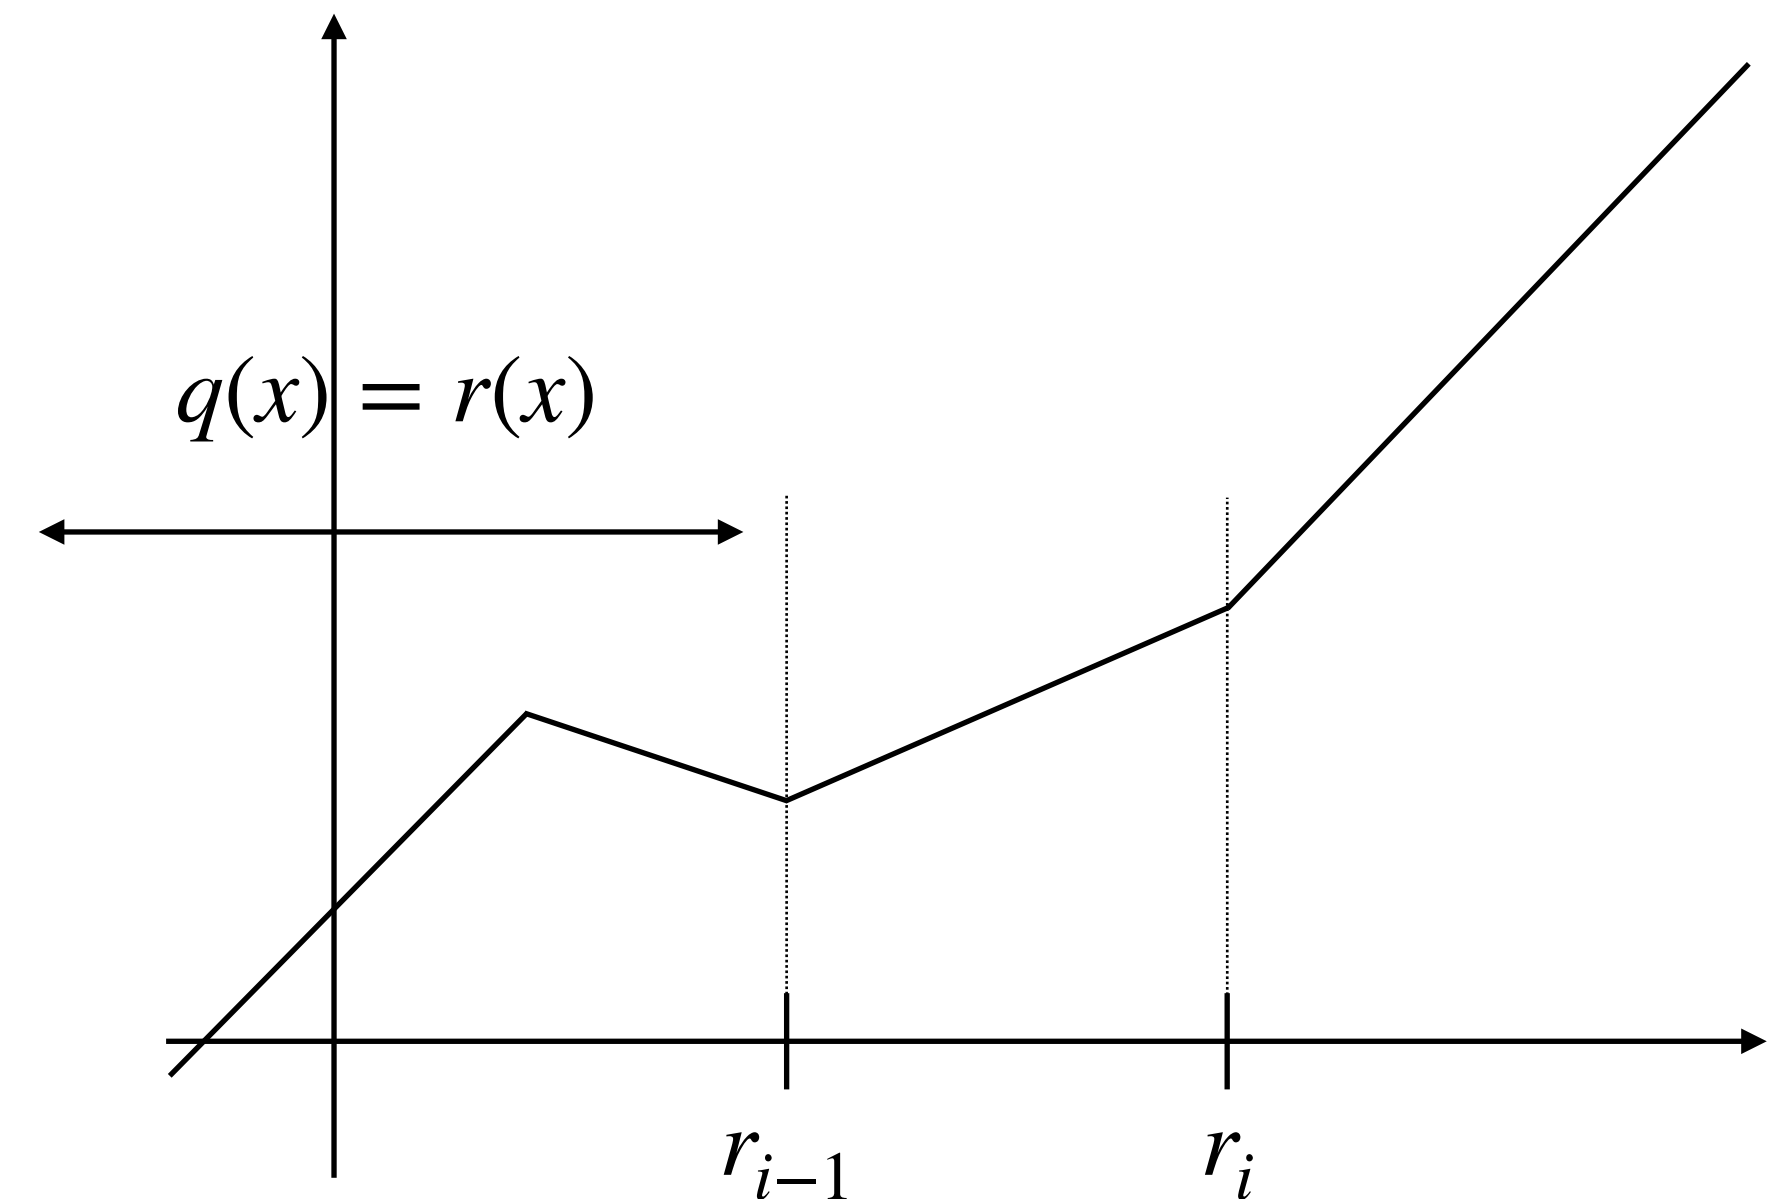
$$\begin{aligned} r(x) &= \tilde{a}_1 x + \tilde{b}_1 + \sum_{j=2}^m \tilde{a}_j (x - \tilde{b}_j)_+ \\ &= \tilde{a}_1 x + \tilde{b}_1 + \sum_{j=2}^i \tilde{a}_j (x - \tilde{b}_j) \\ &= \sum_{j=1}^i \tilde{a}_j x + \tilde{b}_1 - \sum_{j=2}^i \tilde{a}_j \tilde{b}_j \end{aligned}$$

Thus

- $r'(x) = \sum_{j=1}^i \tilde{a}_j = a_i$  good slope
- $r(r_{i-1}) = q(r_{i-1})$  good starting point

$$\implies r(x) = q(x) \text{ for } x \in [r_{i-1}, r_i]$$

Why: two affine functions with the same starting point and the same slope are equal





# Proof by induction - bis

Let's assume that  $r(x) = q(x)$  for  $x \in [0, r_{i-1}]$

For  $x \in [r_{i-1}, r_i]$

$$\begin{aligned} r(x) &= \tilde{a}_1 x + \tilde{b}_1 + \sum_{j=2}^m \tilde{a}_j (x - \tilde{b}_j)_+ \\ &= \underbrace{\tilde{a}_1 x + \tilde{b}_1 + \sum_{j=2}^{i-1} \tilde{a}_j (x - \tilde{b}_j)_+}_{= q(x), x \in [r_{i-2}, r_{i-1}] \text{ by induction}} + \tilde{a}_i (x - \tilde{b}_i) \\ &= a_{i-1} x + b_{i-1} + \tilde{a}_i (x - \tilde{b}_i) \end{aligned}$$

Thus

- $r'(x) = a_{i-1} + \tilde{a}_i = a_i$  good slope
- $r(r_{i-1}) = q(r_{i-1})$  good starting point

$$\implies r(x) = q(x) \text{ for } x \in [r_{i-1}, r_i]$$

Why: two affine functions with the same starting point and the same slope are equal

